

# Introduction to R

Exam Summer Term 2018

*Student Name, Student ID, StudentName@uni-muenster.de*

## Contents

Linear Equation	2
Functions	3
Passenger numbers	4
Porsche	6
Law of large numbers	9
Limits of maxima	10
Student teacher ratio	11
Asymptotic normality	14
Stochastic frontier analysis	16
Variance estimation in GARCH	18

- 
- Answer **7** out of **10** of the following exercises in either German or English.
  - Hand in your solutions before Friday, 13 April 2018 at 10 am.
  - It is advised to regularly check the learnweb and your emails in case of urgent updates.
  - Please send your solutions files to Willi Mutschler. We will confirm the receipt of your work also by email.
  - The solution files should contain your executable and commented script file or preferably a R Notebook.
  - You may use *any* available R package you find fit to solve the exercise.
  - Please label your axes and title in your plots.
  - **All students must work on their own.**

---

## Linear Equation

Solve the linear equation  $A \cdot x = b$  with

$$A = \begin{pmatrix} 1 & 1 & 1 & 3 \\ 2 & 5 & 8 & 6 \\ 4 & 3 & 7 & 9 \\ 3 & 6 & 5 & 1 \end{pmatrix} \text{ and } b = \begin{pmatrix} 1 \\ 4 \\ 3 \\ 5 \end{pmatrix}$$

*Solution:*

```
A = matrix(c(1,2,4,3,1,5,3,6,1,8,7,5,3,6,9,1),4,4)
b = matrix(c(1,4,3,5),4,1)
x=solve(A,b)
print(x)
```

```
##           [,1]
## [1,]  0.20512821
## [2,]  0.76923077
## [3,] -0.05128205
## [4,]  0.02564103
```

and compute the inverse as well as Eigenvalues of  $A$ .

*Solution:*

```
solve(A)
```

```
##           [,1]           [,2]           [,3]           [,4]
## [1,] -0.3589744 -0.43589744  0.38461538  0.23076923
## [2,]  0.6538462  0.11538462 -0.30769231  0.11538462
## [3,] -0.6602564  0.10897436  0.15384615 -0.05769231
## [4,]  0.4551282  0.07051282 -0.07692308 -0.09615385
```

```
eigen(A)$values
```

```
## [1] 17.543241+0.000000i -3.624256+0.000000i  0.040507+1.565859i
## [4]  0.040507-1.565859i
```

---

## Functions

1. Write a function `psum(n,a)` that computes

$$s_{n,a} := \sum_{k=0}^n \frac{k^a}{k^a + 1}$$

for any natural number  $n \in \{1, 2, \dots\}$  and any  $a > 0$ .

*Solution:*

```
psum = function(n,a){  
  if (n==0) 0  
  else ((n^a)/((n^a)+1)+psum(n-1,a))}  
psum(10,10)
```

```
## [1] 9.499006
```

2. Write a function `mymatrix(n)` that returns a  $n \times n$  matrix such that: the first and last row as well as the first and last column contain only ones, whereas the remaining values are zero, e.g.

$$\begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 1 & 1 & 1 & 1 \end{pmatrix}$$

*Solution:*

```
mymatrix = function(n){  
  matrix(c(rep(1,n),rep(c(1,rep(0,n-2),1),n-2),rep(1,n)),n,n)  
  }  
mymatrix(4)
```

```
##      [,1] [,2] [,3] [,4]  
## [1,]    1    1    1    1  
## [2,]    1    0    0    1  
## [3,]    1    0    0    1  
## [4,]    1    1    1    1
```

---

## Passenger numbers

The file **apass.csv** contains monthly data on passenger numbers of US airlines from January 1949 to December 1959.

1. Read the data into a data frame.

*Solution:*

```
apass = read.csv2("../data/apass.csv", header=TRUE)
```

2. Create a vector that contains the corresponding dates. Add this vector to your data frame.


*Solution:*

```
dat1 <- strptime("1949-01-01", "%Y-%m-%d")
dat2 <- strptime("1959-12-01", "%Y-%m-%d")
dates <- seq(dat1, dat2, by = "month")
apass$dates <- dates
```

3. Plot the passenger numbers against the date vector.

*Solution:*

```
plot(apass$dates, apass$y, type = "l")
```




4. Calculate the mean passenger numbers for each month. Plot the means as a bar chart.

*Solution:*

```
strmon <- unique(months(apass$dates))
Z <- rep(NA, 12)
for (i in 1:12) {
  Z[i] <- mean(apass$y[months(apass$dates) == strmon[i]])
}
```

```
}  
barplot(Z, col = "lightblue", names = strmon, main = "Passenger data")
```



---

## Porsche

The file **Porsche911.csv** contains data on Porsche 911 cars, **name** is an id for the owner, **loc** is the location, **age** is the age of the car, **TKM** is the mileage in thousands kilometers and **price** the current price listed on an internet platform for used cars in thousand Euros.

1. Read in the data into a data frame and compute key descriptive statistics (mean, standard deviation, smallest and largest values, quartiles, covariance and correlation) for the variables **age**, **TKM** and **price**.

*Solution:*

```
Porsche911=read.csv2(file=" ../data/Porsche911.csv") [3:5]
summary(Porsche911)
```

```
##           age           TKM           price
##  Min.      : 1.00   Min.      : 21.00   Min.      :11.00
##  1st Qu.:20.25   1st Qu.: 51.95   1st Qu.:22.00
##  Median :29.00   Median : 75.40   Median :29.50
##  Mean    :25.27   Mean    : 70.04   Mean     :37.55
##  3rd Qu.:34.00   3rd Qu.: 93.25   3rd Qu.:39.00
##  Max.    :40.00   Max.    :100.00   Max.     :95.00
```

```
cov(Porsche911)
```

```
##           age           TKM           price
## age      140.3583   194.4894  -251.5397
## TKM      194.4894   686.3065  -451.6083
## price   -251.5397  -451.6083   587.8436
```

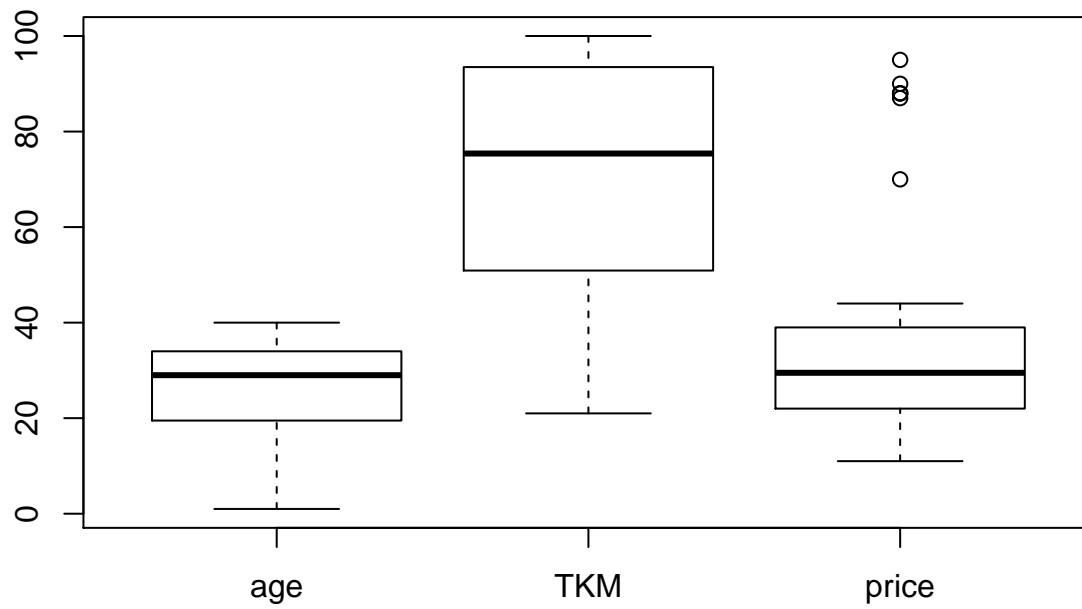
```
cor(Porsche911)
```

```
##           age           TKM           price
## age      1.0000000   0.6266396  -0.8757026
## TKM      0.6266396   1.0000000  -0.7110039
## price   -0.8757026  -0.7110039   1.0000000
```

2. Plot boxplots for **age**, **TKM** and **price**.

*Solution:*


```
boxplot(Porsche911)
```



3. Generate the empirical cumulative distribution function as well as histograms for each of the variables age, TKM and price.

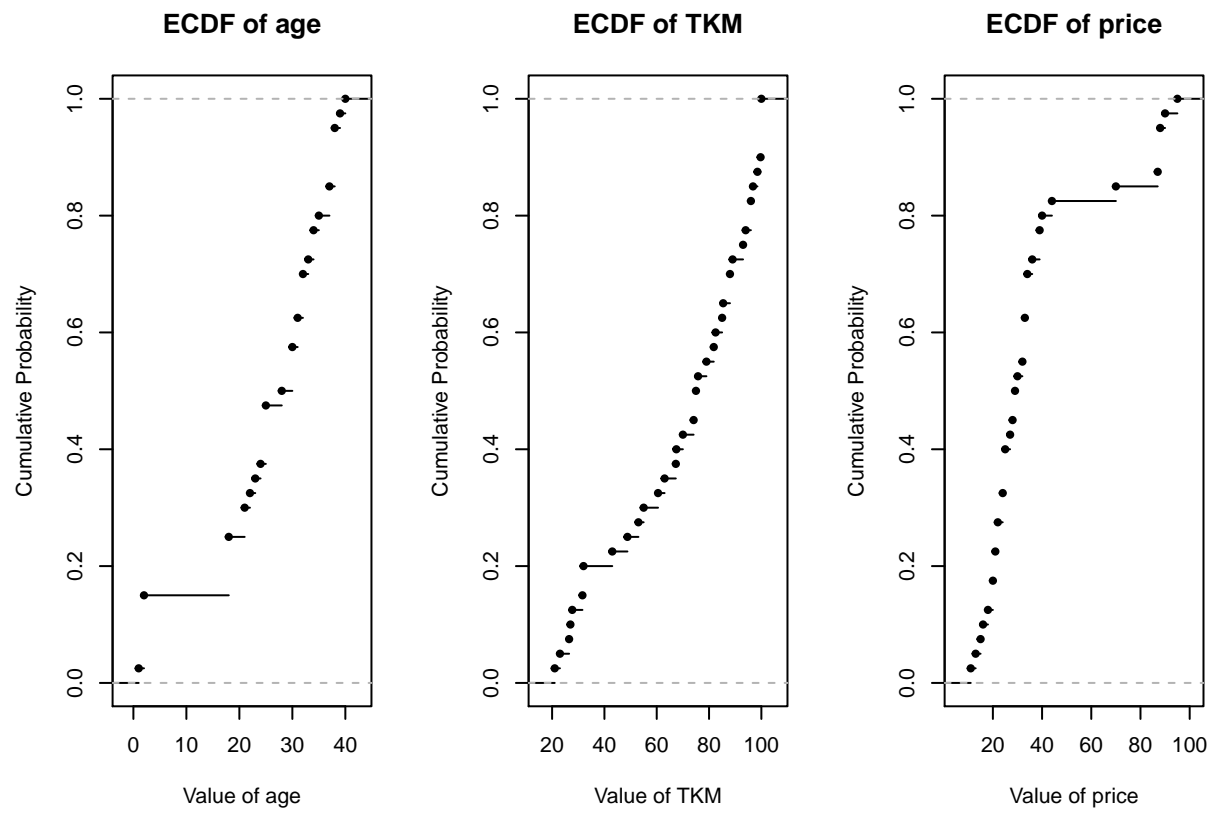
*Solution:*

```
library(MASS)
par(mfrow=c(1,3))
truehist(Porsche911$age)
truehist(Porsche911$TKM)
truehist(Porsche911$price)
```



```
plot(ecdf(Porsche911$age), main = "ECDF of age", xlab = "Value of age", ylab = "Cumulative Probability")
plot(ecdf(Porsche911$TKM), main = "ECDF of TKM", xlab = "Value of TKM", ylab = "Cumulative Probability")
```

```
plot(ecdf(Porsche911$price), main = "ECDF of price", xlab = "Value of price", ylab = "Cumulative Probab
```





---

## Law of large numbers

Let  $X_1, X_2, \dots$  be a sequence  $X_1, X_2, \dots$  from an  $AR(1)$  process:

$$(X_i - \mu) = \rho (X_{i-1} - \mu) + \varepsilon_i$$

where  $\varepsilon_i$  is uniformly distributed on the interval  $[-1, 1]$  and  $|\rho| < 1$ . Define the sequence of random variables

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i.$$


The weak law of large numbers states that the sample average  $\bar{X}_n$  converges in probability towards the expected value  $\mu$  when  $n$  tends to infinity.

Show by simulation that the law of large numbers holds despite the intertemporal dependence in  $X$ . In particular, show the convergence by means of an appropriate plot. Hint: You may set the parameters to e.g.  $\rho = 0.5$  and  $\mu = 2$  (or any other value you find fit.)

**Solution:**

```
n <- 10000
Z <- rep(NA,n)
rho=0.5
mu=2
for (i in 1:n) {
  Z[i] <- mean(filter((1-rho)*mu+runif(i,min=-1,max=1),rho,method="recursive",init=(1-rho)*mu))
}
plot(Z, main="Law of Large numbers for AR(1)");
abline(h=mu,lwd=2,col="red")
```

### Law of Large numbers for AR(1)



---

## Limits of maxima

Let  $X_1, X_2, \dots$  be an i.i.d. sequence of random variables uniformly distributed on the interval  $[0, 1]$ . Define the random variables

$$M_n = \max_{i=1, \dots, n} X_i$$


and its normalized version  $\overline{M}_n = (M_n - 1) \cdot n$ . One can show that the limit distribution of  $\overline{M}_n$  is the Weibull distribution with density  $\exp(x)$ . Write an R program to illustrate that  $\overline{M}_n$  converges in distribution. To this end, set  $n = 100$  and  $R = 1000$  and consider the  $R \times n$  matrix with uniformly distributed random variables  $X_i$ . Show the convergence by means of an appropriate plot.

Note: Try to avoid using loops (use e.g. `apply` instead). You will not get full points if you use a loop.

**Solution:**

```
library(MASS)
n <- 100
R <- 1000 # how many times to draw individual X_i's, note that i = 1,2,...,n
X <- matrix(runif(n*R), R, n) #matrix of N rows and n columns filled with random numbers
Mn <- apply(X, 1, max) # get max of each row
Mnbar <- (Mn-1)*n #standardization
#display the sequence of random variables
truehist(Mnbar, main = paste("n =", toString(n), sep = " "))
coord <- par("usr")
x <- seq(coord[1], coord[2], length.out = 500)
lines(x, exp(x), col = "red")
```

**n = 100**



---

## Student teacher ratio

Load the dataset **caschool.csv** into the object **caschool**. This dataset is discussed in great detail in the textbook of Stock and Watson.

1. Make the following variables accessible:

- test score **testscr**
- student-teacher ratio **str**
- percentage of English language learners **el\_pct**
- expenditures per student **expn\_stu**

*Solution:*

```
caschool <- read.csv("../data/caschool.csv")
testscr <- caschool$testscr
str <- caschool$str
el_pct <- caschool$el_pct
expn_stu <- caschool$expn_stu
```

2. Regress **testscr** on a constant and **str**. Assign the residuals of the regression into the variable **r1**. Now regress **testscr** on an intercept, **str**, **el\_pct** and **expn\_stu**. Put the residuals into the variable **r2**. Compute the sum of squared residuals for both regressions. Display **r1** and **r2** in one plot, where the points of **r2** are marked red.

*Solution:*

```
simple <- lm(testscr~str)
multiple <- lm(testscr~str + el_pct + expn_stu)
r1 <- residuals(simple)
r2 <- residuals(multiple)
sum(r1^2)
```

```
## [1] 144315.5
```


```
sum(r2^2)
```

```
## [1] 85699.69
```

```
sum(r1^2) > sum(r2^2)
```

```
## [1] TRUE
```

```
plot(r1,ylab="")
points(r2,col="red")
```



3. Consider the regression of `testscr` on a constant, `str`, `el_pct` and `expn_stu`. Predict the value of `testscr` for a school district with an average class size (`str`) of 25 students, a percentage of English learners (`el_pct`) of 60% and an average expenditures per student (`expn_stu`) of 4000\$.

*Solution:*

```
predict(multiple, newdata=data.frame(str=25, el_pct=60, expn_stu=4000))
```

```
##          1
## 618.5282
```

4. Reconsider the regression of `testscr` on a constant, `str`, `el_pct` and `expn_stu`. Compute the heteroscedastic robust standard errors.

*Solution:*

```
library(lmtest)
```

```
## Loading required package: zoo
```

```
##
```

```
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      as.Date, as.Date.numeric
```

```
library(sandwich)
```

```
regr <- lm(testscr~str + el_pct + expn_stu)
```

```
summary(regr)
```

```
##
```

```
## Call:
```

```
## lm(formula = testscr ~ str + el_pct + expn_stu)
```

```
##
```

```
## Residuals:
```

```
##      Min      1Q  Median      3Q      Max
## -51.340 -10.111   0.293  10.318  43.181
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 649.577960  15.205716  42.719  < 2e-16 ***
## str          -0.286400   0.480523  -0.596  0.55149
## el_pct        -0.656023   0.039106 -16.776  < 2e-16 ***
## expn_stu       0.003868   0.001412   2.739  0.00643 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 14.35 on 416 degrees of freedom
## Multiple R-squared:  0.4366, Adjusted R-squared:  0.4325
## F-statistic: 107.5 on 3 and 416 DF,  p-value: < 2.2e-16
coeftest(regr, vcov=vcovHC)
```

```
##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 649.5779596  15.6686203  41.4573  < 2e-16 ***
## str          -0.2863998   0.4875129  -0.5875  0.55721
## el_pct        -0.6560228   0.0321143 -20.4278  < 2e-16 ***
## expn_stu       0.0038679   0.0016074   2.4063  0.01655 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

5. Test the null hypothesis that the coefficients on `str` and `expn_stu` both equal 0 and the coefficient on `el_pct` equals  $-0.7$ . Hint: Use the `linearHypothesis` function of the `car` package.

*Solution:*

```
library(car)
linearHypothesis(regr, c("str=0", "expn_stu=0", "el_pct=-.7"))

## Linear hypothesis test
##
## Hypothesis:
## str = 0
## expn_stu = 0
## el_pct = - 0.7
##
## Model 1: restricted model
## Model 2: testscr ~ str + el_pct + expn_stu
##
##      Res.Df    RSS Df Sum of Sq      F    Pr(>F)
## 1         419 89117
## 2         416 85700   3    3416.9 5.5287 0.0009919 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

---

## Asymptotic normality

Consider the multiple linear regression model  $y = X\beta + u$ . In R, generate the matrix  $X$  by executing the following commands:

```
library(MASS)
X <- cbind(1,mvrnorm(n=100,c(5,10),matrix(c(1,0.9,0.9,1),2,2)))
```

Assume that the true coefficient vector is

$$\beta = \begin{pmatrix} 3 \\ 2 \\ -1 \end{pmatrix}$$

and the error terms are i.i.d. uniformly distributed on the interval  $[-1, 1]$ . Hence, the assumption of normally distributed error terms is violated.

1. Write an R program that generates  $R = 10000$  random samples of size  $n = 100$  each. Generate an empty vector  $Z \leftarrow \text{rep}(NA, R)$ . For each sample  $i = 1, \dots, R$ , compute the OLS estimate  $\hat{\beta}$  of  $\beta$  and store the second component of  $\hat{\beta}$  in the  $i$ -th element of the vector  $Z$ .


**Solution:**

```
beta_true <- c(3,2,-1)
n <- 100
R <- 10000
Z <- rep(NA,R)
for (i in 1:R) {
  u <- runif(n,-1,1)
  X <- cbind(1,mvrnorm(n=n,c(5,10),matrix(c(1,0.9,0.9,1),2,2)))
  y <- X%*%beta_true + u
  beta_hat <- solve(t(X)%*%X)%*%t(X)%*%y
  Z[i] <- beta_hat[2]
}
```

2. Plot the histogram of  $Z$ . Compute the mean  $m$  and standard deviation  $s$  of  $Z$  and add the density of  $N(m, s)$  to the plot.

**Solution:**

```
library(MASS)
truehist(Z)
m1 <- mean(Z)
s1 <- sd(Z)
curve(dnorm(x,mean=m1,sd=s1),add=T)
```



---

## Stochastic frontier analysis

Consider the Cobb-Douglas production function

$$y = Ax_1^\alpha x_2^\beta$$

By definition, the production function returns the maximal output for given inputs, and actual production cannot be larger than  $y$ . Due to inefficiencies, actual production could be modeled (in logs) as

$$\ln y = \ln A + \alpha \ln x_1 + \beta \ln x_2 - u$$

where  $u$  is a **non-negative** random variable. Since other disturbances (e.g. measurement errors) can enter the production function, it is more common to add another, **symmetrically distributed**, disturbance term  $v$ ,

$$\ln y = \ln A + \alpha \ln x_1 + \beta \ln x_2 - u + v$$

Assume that  $u$  is exponentially ( $u \sim \text{Exp}(\lambda)$ ) and  $v$  normally ( $v \sim N(0, \sigma^2)$ ) distributed. One can show that if  $u$  and  $v$  are independent then the density function of  $\varepsilon = v - u$  is given by

$$f_\varepsilon(\varepsilon) = \lambda \exp\left(\lambda\varepsilon + \frac{1}{2}\lambda^2\sigma^2\right) \Phi\left(\frac{-\varepsilon}{\sigma} - \lambda\sigma\right)$$

where  $\Phi$  is the distribution function (`pnorm`) of  $N(0, 1)$  and `exp` the exponential function (`exp`).

1. Load the dataset **sfa.csv**. This dataset is an abbreviated version of table F7.2 of Greene, 2008. The original data appeared in Zellner and Revankar, *Generalized Production Functions*, Review of Economic Studies, 36 (1969), 241-250.

**Solution:**

```
sfa <- read.csv("../data/sfa.csv")
head(sfa)
```

```
##      State ValueAdd Capital   Labor
## 1   Alabama  126.148    3.804  31.551
## 2 California 3201.486  185.446 452.844
## 3 Connecticut 690.670   39.712 124.074
## 4   Florida   56.296    6.547  19.181
## 5   Georgia  304.531   11.530  45.534
## 6   Illinois  723.028   58.987  88.391
```

2. Write an R program to estimate the parameters  $A$ ,  $\alpha$ ,  $\beta$ ,  $\lambda$  and  $\sigma$  by maximum likelihood on this dataset.

**Solution:**

```
## negative log-Likelihood function
neg_log_likeli <- function(thet, dat){
  A <- thet[1]
  alpha <- thet[2]
  beta <- thet[3]
  lambda <- thet[4]
  sigma <- thet[5]
  x1 <- dat[,3]
  x2 <- dat[,4]
  y <- dat[,2]
```



```

eps <- log(y) - log(A) - alpha*log(x1) - beta*log(x2)
log_likeli <- sum(log(lambda*exp(lambda*eps+lambda^2/2*sigma^2)*pnorm(-eps/sigma-lambda*sigma,mean=0,
return(-log_likeli)
}

# Estimate parameters
opt <- optim(c(7, 0.3, 0.7, 10, .1), neg_log_likeli, dat=sfa, hessian=T,control=c(maxit=5000))
opt

## $par
## [1] 7.9194127 0.2625018 0.7703623 7.3942188 0.1713962
##
## $value
## [1] -2.860488
##
## $counts
## function gradient
##      512      NA
##
## $convergence
## [1] 0
##
## $message
## NULL
##
## $hessian
##      [,1]      [,2]      [,3]      [,4]      [,5]
## [1,]  9.240673  220.88401  299.17773  1.006997  12.672484
## [2,] 220.884014 6413.44070 8054.50541 22.567215 173.961216
## [3,] 299.177727 8054.50541 10490.92791 31.047041 277.984249
## [4,]  1.006997  22.56722  31.04704  0.222978  -2.947495
## [5,] 12.672484 173.96122  277.98425 -2.947495 914.229088

```

3. Compute the asymptotic standard errors.

**Solution:**

```
sqrt(diag(solve(opt$hessian)))
```

```
## [1] 1.86524162 0.09200213 0.11094965 3.42546237 0.03842999
```

---

## Variance estimation in GARCH

When one considers an iid sample  $X_1, \dots, X_n$  from  $X \sim N(\mu, \sigma^2)$  then one usually estimates the variance  $\sigma^2$  using

$$\hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

The distribution of the normalized estimator for the variance is given by:

$$\frac{(n-1)\hat{\sigma}^2}{\sigma^2} \sim \chi_{n-1}^2$$

where  $\sigma^2$  is the true variance. Consider the case when the observations are not iid but are stochastically dependent over time. To this end, assume that  $X_1, \dots, X_n$  is a time series generated by a  $GARCH(1,1)$  process

$$\begin{aligned} X_i &\sim N(0, \sigma_i^2) \\ \sigma_i^2 &= \omega + \alpha X_{i-1}^2 + \beta \sigma_{i-1}^2 \end{aligned}$$

with  $\omega = 0.1$ ,  $\alpha = 0.1$ ,  $\beta = 0.85$  and sample size equal to  $n = 2500$ . Show by simulations that the distribution of the normalized estimator for the variance is not  $\chi_{n-1}^2$ -distributed. Hint: The true unconditional variance of this GARCH process is  $\sigma^2 = 2$ .

**Solution:**

```
omega <- 0.1
alpha <- 0.1
beta <- 0.85
n <- 2500
R <- 1000

sigma2 <- rep(NA,n)
sigma2[1] <- 0

X <- rep(NA,n)
X[1] <- 0

vargarch <- rep(NA,R)

for (r in 1:R){
  for (i in 2:n){
    sigma2[i] <- omega+alpha*X[i-1]^2+beta*sigma2[i-1]
    X[i] <- rnorm(n=1,mean=0,sd=sqrt(sigma2[i]))
  }
  varest <- 1/(n-1)*sum((X-mean(X))^2)
  vargarch[r] <- (n-1)*varest/2
}

truehist(vargarch,col="lightblue",ylim=c(0,0.006),nbins=30)
g <- seq(1800,3500,length=1000)
lines(g,dchisq(g,df=n-1),lwd=2)
```

