

Introduction to R

Solutions to Exercises

Willi Mutschler

Contents

Introduction

1. Start **R-Studio** and have a look at all menu items.
2. Under Tools - Options choose your preferred setting.
3. Install the packages `dplyr`, `ggplot2`, `tidyr`, `stringr`, `readr`, `foreign`, `readxl`, `haven`, `sandwich`, `prettyR`, `Rcmdr`, `xtable`, `texreg`, and `lmtest`

#you will need to install packages only once

```
install.packages("dplyr")
install.packages("ggplot2")
install.packages("tidyr")
install.packages("stringr")
install.packages("readr")
install.packages("foreign")
install.packages("readxl")
install.packages("haven")
install.packages("sandwich")
install.packages("prettyR")
install.packages("Rcmdr")
install.packages("xtable")
install.packages("texreg")
install.packages("lmtest")
```

More simply: Use the **Packages** Tab of RStudio to install and activate the libraries.

4. Load the `ggplot2` package, list the packages in your memory and detach the `ggplot2` package from the memory

```
library("ggplot2")
search()
```

```
## [1] ".GlobalEnv"      "package:ggplot2"  "package:stats"
## [4] "package:graphics" "package:grDevices" "package:utils"
## [7] "package:datasets" "package:methods"  "Autoloads"
## [10] "package:base"
```

```
detach("package:ggplot2")
search()
```

```
## [1] ".GlobalEnv"      "package:stats"    "package:graphics"
## [4] "package:grDevices" "package:utils"    "package:datasets"
## [7] "package:methods"  "Autoloads"        "package:base"
```

5. The current working directory (where R reads and writes files) can be found by the command `getwd()`. Find your current working directory.

```
getwd()
```

6. Use the command `setwd("c:/path")` or `setwd(choose.dir())` (only available on windows) to change the working directory to drive *c:* and path */path*. Note that the path name is structured by slashes (/), **not** backslashes (\). Change the working directory to *c:/temp* and check if the change has been successful.

```
setwd("C:/temp")
#setwd(choose.dir()) # on windows
getwd()
```

Hint: The working directory can also be changed via the menu: *Session – Set Working Directory*.

7. Open a new script file. Type the commands to perform the following assignments:

$$a = \frac{3 \cdot (4 + 9)}{8 - 12.5}$$
$$b = (1, 4, 1999, 2011)$$
$$d = 2\pi$$
$$e = a + d$$

Save the script and quit R.

```
a <- 3 * (4 + 9) / (8 - 12.5)
b <- c(1, 4, 1999, 2011)
d <- 2 * pi
e <- a + d
```

8. Start R and re-open the script. Mark all lines (*Ctrl-A*) and execute them. Print *a, b, d, e*.

```
## [1] -8.666667
```

```
## [1]      1      4 1999 2011
```

```
## [1] 6.283185
```

```
## [1] -2.383481
```

9. Why is *c* not used as a variable?

Because c() is the built-in concatenation command! Do not use c for variable or function declaration.

Logical Operators

1. Use the command `c()` to define the vectors

$$x = (-1, 0, 1, 4, 9, 2, 1, 4.5, 1.1, -0.9)$$
$$y = (1, 1, 2, 2, 3, 3, 4, 4, 5, NA)$$

```
x <- c(-1, 0, 1, 4, 9, 2, 1, 4.5, 1.1, -0.9)
y <- c(1, 1, 2, 2, 3, 3, 4, 4, 5, NA)
```

2. Determine the lengths of the vectors using `length()` and check if `length(x)==length(y)`.

```
length(x); length(y); length(x) == length(y)
```

```
## [1] 10
```

```
## [1] 10
```

```
## [1] TRUE
```

3. Perform the following logical operations:

$$\begin{aligned}x &< y \\x &< 0 \\x + 3 &\geq 0 \\y &< 0 \\x &< 0 \text{ or } y < 0\end{aligned}$$

```
x < y
```

```
## [1] TRUE TRUE TRUE FALSE FALSE TRUE TRUE FALSE TRUE NA
```

```
x < 0
```

```
## [1] TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE
```

```
x + 3 >= 0
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

```
y < 0
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE NA
```

```
x < 0 | y < 0
```

```
## [1] TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE
```

4. Use `all` to check if all elements of $x + 3 \geq 0$.

```
all(x + 3 >= 0)
```

```
## [1] TRUE
```

5. Use `all` to check if all elements of $y > 0$. Use `any` to check if at least one element of $y > 0$.

```
all(y > 0)
```

```
## [1] NA
```

```
all(y > 0, na.rm = TRUE)
```

```
## [1] TRUE
```

```
any(y > 0)
```

```
## [1] TRUE
```

Arithmetic operators and mathematical functions

1. Define the vectors

$$x = (-1, 0, 1, 4, 9, 2, 1, 4.5, 1.1, -0.9)$$

$$y = (1, 1, 2, 2, 3, 3, 4, 4, 5, NA).$$

and compute $x + y$ and xy and y/x .

```
x <- c(-1, 0, 1, 4, 9, 2, 1, 4.5, 1.1, -0.9)
y <- c(1, 1, 2, 2, 3, 3, 4, 4, 5, NA)

x + y

## [1] 0.0 1.0 3.0 6.0 12.0 5.0 5.0 8.5 6.1 NA
x * y

## [1] -1.0 0.0 2.0 8.0 27.0 6.0 4.0 18.0 5.5 NA
y / x

## [1] -1.0000000      Inf 2.0000000 0.5000000 0.3333333 1.5000000
## [7] 4.0000000 0.8888889 4.5454545      NA
```

2. Compute $\ln(x)$. Determine the length of the result vector.

```
log(x)

## Warning in log(x): NaNs wurden erzeugt

## [1]      NaN      -Inf 0.00000000 1.38629436 2.19722458 0.69314718
## [7] 0.00000000 1.50407740 0.09531018      NaN

length(log(x))

## Warning in log(x): NaNs wurden erzeugt

## [1] 10
```

3. Use `any` to check if the vector x contains elements satisfying $\sqrt{x} \geq 2$.

```
any(sqrt(x) >= 2)

## Warning in sqrt(x): NaNs wurden erzeugt

## [1] TRUE
```

4. Compute

$$a = \sum_{i=1}^{10} x_i$$

$$b = \sum_{i=1}^{10} y_i^2.$$

Use the `na.rm=TRUE` option (**na-remove**) of the `sum` command to drop the `NA` in y .

```
a <- sum(x)
print(a)

## [1] 20.7
```

```
b <- sum(y^2, na.rm = TRUE)
print(b)
```

```
## [1] 85
```

5. Compute

$$\sum_{i=1}^{10} x_i y_i^2.$$

```
sum(x * (y^2), na.rm = TRUE)
```

```
## [1] 233.5
```

6. The `sum` command is a convenient way to count the number of elements satisfying a certain condition. Count the number of elements of $x > 0$.

```
sum(x > 0)
```

```
## [1] 7
```

```
sum(!is.na(y))
```

```
## [1] 9
```

7. Predict what the following commands will return:

```
x^y
```

```
## [1] -1.00000 0.00000 1.00000 16.00000 729.00000 8.00000 1.00000
## [8] 410.06250 1.61051 NA
```

```
x^(1/y)
```

```
## [1] -1.000000 0.000000 1.000000 2.000000 2.080084 1.259921 1.000000
## [8] 1.456475 1.019245 NA
```

```
log(exp(y))
```

```
## [1] 1 1 2 2 3 3 4 4 5 NA
```

```
y*c(-1,1)
```

```
## [1] -1 1 -2 2 -3 3 -4 4 -5 NA
```

```
x+c(-1,0,1)
```

```
## Warning in x + c(-1, 0, 1): Länge des längeren Objektes
##      ist kein Vielfaches der Länge des kürzeren Objektes
## [1] -2.0 0.0 2.0 3.0 9.0 3.0 0.0 4.5 2.1 -1.9
```

```
sum(y*c(-1,1),na.rm=TRUE)
```

```
## [1] -5
```

Matrix functions

1. Define the matrix

$$X = \begin{bmatrix} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 9 \end{bmatrix},$$

print its transpose, its dimensions and its determinant.

```
X <- matrix(c(1,2,3,4,5,6,7,8,9), nrow = 3, ncol = 3)
t(X)
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    3
## [2,]    4    5    6
## [3,]    7    8    9
```

```
dim(X)
```

```
## [1] 3 3
```

```
det(X)
```

```
## [1] 0
```

2. Compute the trace of X (i.e. the sum of its diagonal elements).

```
sum(diag(X))
```

```
## [1] 15
```

3. Type `diag(X) <- c(7,8,9)` to change the diagonal elements. Compute the eigenvalues of (the new) X .

```
diag(X) <- c(7, 8, 9)
eigen(X)
```

```
## eigen() decomposition
## $values
## [1] 18.000000  4.732051  1.267949
##
## $vectors
##      [,1]      [,2]      [,3]
## [1,] -0.5773503 -0.90894503 -0.3239853
## [2,] -0.5773503  0.41277422 -0.6824097
## [3,] -0.5773503  0.05862061  0.6552484
```

Is X positive definite?

```
# Yes, since all eigenvalues are positive.
```

4. Invert X and compute the eigenvalues of X^{-1} .

```
solve(X) # inverts X
```

```
##      [,1]      [,2]      [,3]
## [1,] 0.22222222 0.05555556 -0.2222222
## [2,] 0.05555556 0.38888889 -0.3888889
## [3,] -0.11111111 -0.27777778 0.4444444
```

```
eigen(solve(X))
```

```
## eigen() decomposition
## $values
```

```
## [1] 0.78867513 0.21132487 0.05555556
##
## $vectors
##      [,1]      [,2]      [,3]
## [1,] -0.3239853 -0.90894503 0.5773503
## [2,] -0.6824097  0.41277422 0.5773503
## [3,]  0.6552484  0.05862061 0.5773503
```

5. Define the vector $a = (1, 3, 2)$ and compute $a \cdot X$, $a \%*\% X$, and $X \%*\% a$.

```
a <- c(1, 3, 2)
a * X
```

```
##      [,1] [,2] [,3]
## [1,]    7    4    7
## [2,]    6   24   24
## [3,]    6   12   18
```

```
a \%*\% X
```

```
##      [,1] [,2] [,3]
## [1,]   19   40   49
```

```
X \%*\% a
```

```
##      [,1]
## [1,]   33
## [2,]   42
## [3,]   39
```

6. Compute the quadratic form $a' X a$.

```
t(a) \%*\% X \%*\% a
```

```
##      [,1]
## [1,]  237
```

7. Define the matrices $I = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$, $Y = \begin{bmatrix} 1 & 4 & 7 & 1 & 0 & 0 \\ 2 & 5 & 8 & 0 & 1 & 0 \\ 3 & 6 & 9 & 0 & 0 & 1 \end{bmatrix}$ and $Z = \begin{bmatrix} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 9 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$.

```
I <- diag(3)
Y <- cbind(matrix(1:9,3,3),I)
Z <- rbind(matrix(1:9,3,3),I)
```

8. Predict what the following commands will return:

```
cbind(1,X)
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    7    4    7
## [2,]    1    2    8    8
## [3,]    1    3    6    9
```

```
rbind(Y,c(1,2,3))
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]
## [1,]    1    4    7    1    0    0
## [2,]    2    5    8    0    1    0
```

```
## [3,] 3 6 9 0 0 1
## [4,] 1 2 3 1 2 3

X%*%I

##      [,1] [,2] [,3]
## [1,] 7 4 7
## [2,] 2 8 8
## [3,] 3 6 9

dim(X%*%Y)

## [1] 3 6

t(Y)+Z

##      [,1] [,2] [,3]
## [1,] 2 6 10
## [2,] 6 10 14
## [3,] 10 14 18
## [4,] 2 0 0
## [5,] 0 2 0
## [6,] 0 0 2

solve(t(Z)%*%Z)%*%(t(Z)%*%Z)

##      [,1] [,2] [,3]
## [1,] 1.000000e+00 -1.265481e-14 -2.042897e-14
## [2,] 2.331468e-15 1.000000e+00 -2.275957e-15
## [3,] -1.221245e-15 -7.771561e-16 1.000000e+00
```

Set operations and special functions

1. Define the vectors

$$x = (-1, 0, 1, 4, 9, 2, 1, 4.5, 1.1, -0.9)$$

$$y = (1, 1, 2, 2, 3, 3, 4, 4, 5, NA).$$

and compute $x \cup y$. Determine the lengths of x , y and $x \cup y$.

```
x <- c(-1, 0, 1, 4, 9, 2, 1, 4.5, 1.1, -0.9)
y <- c(1, 1, 2, 2, 3, 3, 4, 4, 5, NA)
union(x, y)

## [1] -1.0 0.0 1.0 4.0 9.0 2.0 4.5 1.1 -0.9 3.0 5.0 NA

length(x)

## [1] 10

length(y)

## [1] 10

length(union(x, y))

## [1] 12
```


2. Count the number of elements of y that are element of x .

```
sum(unique(y) %in% x)
```

```
## [1] 3
```

```
length(intersect(y,x))
```

```
## [1] 3
```

3. Determine the length of the vector of unique elements of y .

```
length(unique(y))
```

```
## [1] 6
```

4. Compute the vector z with elements

$$z_i = \sum_{j=1}^i x_j$$

for $i = 1, \dots, 10$.

```
z <- cumsum(x)
print(z)
```

```
## [1] -1.0 -1.0  0.0  4.0 13.0 15.0 16.0 20.5 21.6 20.7
```

5. Find the position of the largest element of x .

```
which.max(x)
```

```
## [1] 5
```

Sequences and replications

1. Generate the vectors

$$\begin{aligned}x_1 &= (1, 2, 3, \dots, 9) \\x_2 &= (0, 1, 0, 1, 0, 1, 0, 1) \\x_3 &= (1, 1, 1, 1, 1, 1, 1, 1) \\x_4 &= (-1, 1, -1, 1, -1, 1) \\x_5 &= (1980, 1985, 1990, \dots, 2010) \\x_6 &= (0, 0.01, 0.02, \dots, 0.99, 1)\end{aligned}$$

```
x1 <- 1:9 # or c(1:9)
x2 <- rep(c(0, 1), times = 4)
x3 <- rep(1, times = 8)
x4 <- rep(c(-1, 1), times = 3)
x5 <- seq(from = 1980, to = 2010, by = 5)
x6 <- seq(0, 1, by = 0.01)
```

2. Replications can also be generated for vectors of strings (characters). Type

```
a <- c("a", "b", "c")
rep(a, 3)
```

```
## [1] "a" "b" "c" "a" "b" "c" "a" "b" "c"
```

```
rep(a, times = 3)
```

```
## [1] "a" "b" "c" "a" "b" "c" "a" "b" "c"
```

```
rep(a, each = 3)
```

```
## [1] "a" "a" "a" "b" "b" "b" "c" "c" "c"
```

3. Generate a grid of $n = 500$ equidistant points on the interval $[-\pi, \pi]$.

```
seq(-pi, pi, length = 500)
```

4. Compare $1:10+1$ and $1:(10+1)$.

```
1:10 + 1      # sequence from 2 to 11
```

```
## [1]  2  3  4  5  6  7  8  9 10 11
```

```
1:(10 + 1)    # sequence from 1 to 11
```

```
## [1]  1  2  3  4  5  6  7  8  9 10 11
```

5. Predict what the following commands will return:

```
rep("bla",10)
```

```
## [1] "bla" "bla" "bla" "bla" "bla" "bla" "bla" "bla" "bla" "bla"
```

```
rep(rep(1:3,2),each=4)
```

```
## [1] 1 1 1 1 2 2 2 2 3 3 3 3 1 1 1 1 2 2 2 2 3 3 3 3
```

```
rep(c(1,6,NA,2),times=c(2,2,5,3))
```

```
## [1]  1  1  6  6 NA NA NA NA NA  2  2  2
```

Reading and writing text files

Consider the files **bsp1.txt**, **bsp2.txt** and **bsp3.txt**. The three files contain computer generated random numbers.

1. Read the file **bsp1.txt** into a data frame **bsp1**. Have a look at the file and the data format *before* you decide which reading command you use (`read.csv`, `read.csv2` or `read.table`). Print the data frame. If the data frame is too large for your screen, you can use the commands `head` and `tail` to print just parts of it.

```
bsp1 <- read.csv2("data/bsp1.txt") # German format (sep=';', dec=','), therefore read.csv2
print(bsp1)
```

```
##      Variable1 Variable2 Variable3 Variable4
## 1           15         12          4        1.12
## 2              0          0         14        0.23
## 3           11          3          3        0.12
## 4              0          1         12        1.94
```

## 5	0	1	10	0.66
## 6	2	6	11	0.54
## 7	6	10	6	0.07
## 8	14	2	2	0.75
## 9	0	0	2	1.41
## 10	3	6	3	2.11
## 11	0	11	1	0.83
## 12	1	5	10	0.11
## 13	10	0	8	0.74
## 14	4	2	1	0.65
## 15	1	1	2	0.13
## 16	1	1	3	1.03
## 17	0	5	3	2.35
## 18	3	17	6	0.27
## 19	2	14	10	0.02
## 20	9	8	4	0.02
## 21	7	2	21	0.22
## 22	15	0	4	0.21
## 23	1	0	7	0.73
## 24	0	8	13	0.08
## 25	0	1	0	1.15
## 26	0	0	0	0.34
## 27	0	9	4	1.66
## 28	10	14	1	0.92
## 29	3	3	3	1.69
## 30	0	24	3	0.06
## 31	0	1	1	0.40
## 32	5	4	16	0.10
## 33	0	0	0	0.82
## 34	12	0	1	1.49
## 35	7	8	18	0.74
## 36	7	10	6	0.44
## 37	19	3	26	0.55
## 38	1	13	4	0.74
## 39	1	1	0	0.49
## 40	5	5	1	0.71
## 41	1	1	0	0.21
## 42	3	1	2	1.51
## 43	12	9	1	0.42
## 44	7	1	0	0.67
## 45	1	5	5	2.19
## 46	7	5	0	0.93
## 47	3	3	2	0.35
## 48	1	1	11	1.28
## 49	6	4	2	3.01
## 50	1	6	21	0.47
## 51	4	2	2	0.86
## 52	15	2	2	0.03
## 53	15	13	2	0.38
## 54	3	1	0	0.28
## 55	2	1	0	1.24
## 56	8	17	0	2.26
## 57	3	7	3	1.11
## 58	2	7	2	0.48

## 59	0	1	8	0.48
## 60	7	7	0	0.36
## 61	5	51	0	1.06
## 62	11	0	14	0.20
## 63	5	0	12	2.03
## 64	5	1	23	0.32
## 65	1	4	2	0.51
## 66	3	4	1	0.32
## 67	2	1	0	0.78
## 68	5	2	13	1.48
## 69	1	1	1	0.67
## 70	2	0	0	0.06
## 71	2	1	6	0.51
## 72	2	3	15	0.64
## 73	3	1	7	0.55
## 74	8	4	7	0.03
## 75	3	10	3	1.65
## 76	0	11	6	0.80
## 77	4	1	7	1.18
## 78	11	0	7	0.61
## 79	3	1	7	0.97
## 80	6	2	1	1.79
## 81	11	10	1	0.81
## 82	3	1	1	0.11
## 83	0	27	10	1.07
## 84	2	10	2	0.85
## 85	0	15	4	1.77
## 86	36	0	2	0.98
## 87	1	4	1	1.12
## 88	0	12	8	0.79
## 89	4	2	11	1.01
## 90	8	11	0	0.43
## 91	9	0	1	0.24
## 92	23	17	2	2.58
## 93	0	4	6	0.41
## 94	5	11	1	0.46
## 95	9	14	4	2.08
## 96	4	3	0	0.41
## 97	6	8	0	1.20
## 98	4	5	24	0.56
## 99	12	27	6	1.77
## 100	6	15	4	1.26

```
head(bsp1)
```

##	Variable1	Variable2	Variable3	Variable4
## 1	15	12	4	1.12
## 2	0	0	14	0.23
## 3	11	3	3	0.12
## 4	0	1	12	1.94
## 5	0	1	10	0.66
## 6	2	6	11	0.54

```
tail(bsp1)
```

```
##      Variable1 Variable2 Variable3 Variable4
## 95           9         14          4        2.08
## 96           4          3          0        0.41
## 97           6          8          0        1.20
## 98           4          5         24        0.56
## 99          12         27          6        1.77
## 100          6         15          4        1.26
```

2. Read the files **bsp2.txt** and **bsp3.txt** into data frames **bsp2** and **bsp3**. Note that **bsp2.txt** contains both numeric and character entries. It is usually advisable to set the option **as.is=TRUE** when reading characters (strings).

```
# international format (sep=',', dec='.'), therefore read.csv:
bsp2 <- read.csv("data/bsp2.txt")
# See help file for default settings
bsp3 <- read.table("data/bsp3.txt", dec = ".", sep = ",")
```

3. Print the class of **bsp2**, its dimension, and its variable names (use **names**).

```
class(bsp2)
```

```
## [1] "data.frame"
```

```
dim(bsp2)
```

```
## [1] 40  5
```

```
names(bsp2)
```

```
## [1] "X" "Y" "Z" "U" "V"
```

4. Print a summary of **bsp3**.

```
summary(bsp3)
```

```
##      V1      V2      V3      V4
## Min.   :-5.3800 Min.   :-8.0800 Min.   :-5.814 Min.   :-5.527
## 1st Qu.: -0.2288 1st Qu.: -2.2403 1st Qu.: -1.028 1st Qu.: -0.476
## Median : 2.8390 Median : 0.2575 Median : 1.657 Median : 2.219
## Mean   : 2.3281 Mean   : 0.1810 Mean   : 1.520 Mean   : 1.911
## 3rd Qu.: 4.6227 3rd Qu.: 2.7015 3rd Qu.: 3.517 3rd Qu.: 3.901
## Max.    : 9.2180 Max.    : 8.4800 Max.    : 8.773 Max.    : 9.045
```

5. Compute the mean and the standard deviation of each column of **bsp3**.

```
apply(bsp3, 2, sd)
```

```
##      V1      V2      V3      V4
## 3.230965 3.775867 3.422948 3.345753
```

```
apply(bsp3, 2, mean)
```

```
##      V1      V2      V3      V4
## 2.32806 0.18104 1.51952 1.91102
```

6. Create a small data frame **a** with two variables

x	y
1	4

x	y
2	5
3	6

and write it to a file `smalldataframe.csv` in your working directory.

```
x <- 1:3
y <- 4:6
smalldataframe <- data.frame(x, y)
write.csv2(smalldataframe, file = "data/smalldataframe.csv") # write.csv2 for German Excel
```

7. Read the (large) file `lest2001.csv` into a data frame `x`. The file is the campus file of the German income tax records 2001 (the data are provided by the Research Data Centre of the Federal Statistical Office, they are described in `lest2001.pdf`). Take care to set the options of the `read.csv` or `read.table` command correctly. The data format is as follows:

- All data entries are separated by semi-colons.
- The first row contains the variables names.
- Missing values are denoted by a dot.
- Apart from the last column all data are integer values.
- The decimal sign in the last column is a dot.

Execute `y <- x$zve`. The variable `y` now contains the taxable income (zu versteuerndes Einkommen). Compute its range, its median, its mean, its variance, and the 0.01- and 0.99-quantiles. Remember to include the option `na.rm=TRUE` in the functions.

```
lest2001 <- read.table("data/lest2001.csv", header = TRUE, na.strings = ".", dec = ".", sep = ";")
head(lest2001)
```

```
##      ef0 ef8 ef11 ef12 ef48 ab merker tabelle alter_a alter_b kinder region
## 1      1      0      2      2      4      1      1      2      4      4      1      1
## 2      2      1     NA      2      1      1      1      1      0      5      0      2
## 3      3      0      2     NA      3      1      1      2      4      4      1      1
## 4      4      0      6     NA      0      1      1      1      5      0      0      1
## 5      5      0      2     NA      1      1      1      1      3      0      0      1
## 6      6      0      2      2      4      1      1      2      3      3      1      2
##      sde   gde einkommen   zve est_tarif est_fest c65101 c65102 c65121
## 1 58258 58258   45239 45239   8044   7532      NA      NA      NA
## 2 17536 17536   15479 15479   1985   1985      NA      NA      NA
## 3 29024 29024   23573 23573   2087   2087      NA      NA      NA
## 4 10162 10162    5398  5398      0      0      NA      NA      1
## 5 68216 68216   66009 66009  22143  22143      NA      NA      NA
## 6 54512 54512   44106 44106   7691   7180      NA      NA      NA
##      c65122 c65141 c65142 c65161 c65162 c65221 c65222 c65241 c65242 c65261
## 1      NA      NA      NA      1      1      NA      NA      1      NA      NA
## 2      NA      NA      NA      NA      1      NA      NA      NA      NA      NA
## 3      NA      NA      NA      1      NA      NA      NA      NA      NA      NA
## 4      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
## 5      NA      NA      NA      1      NA      NA      NA      NA      NA      NA
## 6      NA      NA      NA      1      1      NA      NA      1      NA      NA
##      c65262 c65413 c65427 c65479 c65490 c65727 samplingweight
## 1      NA      1      NA      1      NA      NA      111.1101
## 2      NA      NA      1      NA      NA      NA      111.0959
## 3      NA      1      1      NA      NA      NA      111.1093
## 4      NA      1      1      NA      NA      NA      111.0878
```

```
## 5      NA      1      1      NA      NA      NA      111.1004
## 6      NA      1      1      1      NA      NA      111.0951
```

```
y <- lest2001$zve
range(y, na.rm = TRUE)
```

```
## [1] -509791 44887210
```

```
median(y, na.rm = TRUE)
```

```
## [1] 22457
```

```
var(y, na.rm = TRUE)
```

```
## [1] 296245129137
```

```
quantile(y, p = c(0.01, 0.99), na.rm = TRUE)
```

```
##          1%          99%
## -5247.16 1282557.64
```

- Import the dataset **swimming_pools.csv**. It contains data on swimming pools in Brisbane, Australia (Source: data.gov.au). The file contains the column names in the first row and uses a comma to separate values within rows.

```
pools <- read.csv("data/swimming_pools.csv", stringsAsFactors = FALSE)
str(pools) # Check the structure of pools
```

```
## 'data.frame':   20 obs. of  4 variables:
## $ Name      : chr  "Acacia Ridge Leisure Centre" "Bellbowrie Pool" "Carole Park" "Centenary Pool (in
## $ Address   : chr  "1391 Beaudesert Road, Acacia Ridge" "Sugarwood Street, Bellbowrie" "Cnr Boundary
## $ Latitude  : num  -27.6 -27.6 -27.6 -27.5 -27.4 ...
## $ Longitude: num  153 153 153 153 153 ...
```

- Import **hotdogs.txt**, containing information on sodium and calorie levels in different hotdogs (Source: UCLA). The dataset has 3 variables (type, calories and sodium), but the variable names are not available in the first line of the file. The file uses tabs as field separators.

```
hotdogs <- read.delim("data/hotdogs.txt", header = FALSE)
summary(hotdogs) # Summarize hotdogs
```

```
##          V1          V2          V3
## Beef      :20   Min.   : 86.0   Min.   :144.0
## Meat      :17   1st Qu.:132.0   1st Qu.:362.5
## Poultry:17   Median :145.0   Median :405.0
##           Mean    :145.4   Mean    :424.8
##           3rd Qu.:172.8   3rd Qu.:503.5
##           Max.    :195.0   Max.    :645.0
```

```
hotdogs <- read.delim("data/hotdogs.txt", header = FALSE, col.names = c("type", "calories", "sodium"))
str(hotdogs)
```

```
## 'data.frame':   54 obs. of  3 variables:
## $ type      : Factor w/ 3 levels "Beef","Meat",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ calories: int  186 181 176 149 184 190 158 139 175 148 ...
## $ sodium   : int  495 477 425 322 482 587 370 322 479 375 ...
```

```
# Edit the colClasses argument to import the data correctly: hotdogs2
hotdogs2 <- read.delim("data/hotdogs.txt", header = FALSE,
                      col.names = c("type", "calories", "sodium"),
```

```
colClasses = c("factor", "NULL", "numeric"))
str(hotdogs2)

## 'data.frame':   54 obs. of  2 variables:
## $ type   : Factor w/ 3 levels "Beef","Meat",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ sodium: num  495 477 425 322 482 587 370 322 479 375 ...
```

Import data using readr

Use the package `readr` to import the datasets used in the previous exercise, i.e. `bsp1.txt`, `bsp2.txt`, `bsp3.txt`, `lest2001.csv`, `swimming_pools.csv` and `hotdogs.txt`.

```
library(readr)
```

Import Excel data

1. Load the `readxl` package.

```
library(readxl)
```

2. Print the names of all worksheets in the excel file `urbanpop.xlsx`. This dataset is a subset of the `gapminder` dataset.

```
excel_sheets("data/urbanpop.xlsx")
```

```
## [1] "1960-1966" "1967-1974" "1975-2011"
```

3. Now read the sheets, one by one, using `read_excel` and put these into a list.

```
pop_1 <- read_excel("data/urbanpop.xlsx", sheet = 1)
pop_2 <- read_excel("data/urbanpop.xlsx", sheet = 2)
pop_3 <- read_excel("data/urbanpop.xlsx", sheet = 3)
pop_list <- list(pop_1, pop_2, pop_3)
pop_list
```

```
## [[1]]
## # A tibble: 209 x 8
##   country      `1960`    `1961`    `1962`    `1963`    `1964`    `1965`    `1966`
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 Afghanistan 769308  8.15e5  8.59e5  9.04e5  9.51e5  1.00e6  1.06e6
## 2 Albania      494443  5.12e5  5.29e5  5.47e5  5.66e5  5.84e5  6.03e5
## 3 Algeria      3293999  3.52e6  3.74e6  3.97e6  4.22e6  4.49e6  4.65e6
## 4 American Sam~      NA  1.37e4  1.42e4  1.48e4  1.54e4  1.60e4  1.67e4
## 5 Andorra       NA  8.72e3  9.70e3  1.07e4  1.19e4  1.31e4  1.42e4
## 6 Angola        521205  5.48e5  5.80e5  6.12e5  6.45e5  6.79e5  7.18e5
## 7 Antigua and ~  21699  2.16e4  2.17e4  2.17e4  2.18e4  2.19e4  2.20e4
## 8 Argentina     15224096  1.55e7  1.59e7  1.63e7  1.67e7  1.70e7  1.74e7
## 9 Armenia       957974  1.01e6  1.06e6  1.12e6  1.17e6  1.23e6  1.28e6
## 10 Aruba        24996  2.81e4  2.85e4  2.88e4  2.89e4  2.91e4  2.93e4
## # ... with 199 more rows
```



```
##
## [[2]]
## # A tibble: 209 x 9
##   country    `1967`  `1968`  `1969`  `1970`  `1971`  `1972`  `1973`  `1974`
##   <chr>      <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 Afghanist~ 1.12e6  1.18e6  1.25e6  1.32e6  1.41e6  1.50e6  1.60e6  1.70e6
## 2 Albania    6.21e5  6.40e5  6.59e5  6.78e5  6.99e5  7.20e5  7.42e5  7.63e5
## 3 Algeria    4.83e6  5.02e6  5.22e6  5.43e6  5.62e6  5.82e6  6.02e6  6.24e6
## 4 American ~ 1.73e4  1.80e4  1.86e4  1.92e4  1.98e4  2.03e4  2.07e4  2.12e4
## 5 Andorra    1.54e4  1.67e4  1.81e4  1.95e4  2.09e4  2.24e4  2.39e4  2.55e4
## 6 Angola     7.57e5  7.98e5  8.41e5  8.86e5  9.55e5  1.03e6  1.10e6  1.18e6
## 7 Antigua a~ 2.21e4  2.21e4  2.22e4  2.22e4  2.26e4  2.29e4  2.32e4  2.35e4
## 8 Argentina  1.78e7  1.81e7  1.85e7  1.89e7  1.93e7  1.98e7  2.02e7  2.07e7
## 9 Armenia    1.34e6  1.39e6  1.45e6  1.51e6  1.56e6  1.62e6  1.68e6  1.74e6
## 10 Aruba      2.94e4  2.96e4  2.97e4  2.99e4  3.01e4  3.03e4  3.05e4  3.06e4
## # ... with 199 more rows
##
## [[3]]
## # A tibble: 209 x 38
##   country `1975`  `1976`  `1977`  `1978`  `1979`  `1980`  `1981`  `1982`  `1983`
##   <chr>    <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 Afghan~ 1.79e6  1.91e6  2.02e6  2.14e6  2.27e6  2.40e6  2.49e6  2.59e6  2.69e6
## 2 Albania 7.85e5  8.08e5  8.31e5  8.54e5  8.78e5  9.02e5  9.27e5  9.52e5  9.78e5
## 3 Algeria 6.46e6  6.77e6  7.10e6  7.45e6  7.81e6  8.19e6  8.64e6  9.11e6  9.59e6
## 4 Americ~ 2.16e4  2.20e4  2.25e4  2.29e4  2.35e4  2.42e4  2.52e4  2.63e4  2.77e4
## 5 Andorra 2.70e4  2.84e4  2.97e4  3.10e4  3.26e4  3.44e4  3.64e4  3.86e4  4.10e4
## 6 Angola  1.27e6  1.37e6  1.48e6  1.60e6  1.72e6  1.86e6  2.02e6  2.19e6  2.37e6
## 7 Antigu~ 2.38e4  2.40e4  2.42e4  2.43e4  2.44e4  2.43e4  2.42e4  2.39e4  2.36e4
## 8 Argent~ 2.11e7  2.16e7  2.20e7  2.24e7  2.29e7  2.33e7  2.38e7  2.43e7  2.48e7
## 9 Armenia 1.80e6  1.85e6  1.90e6  1.95e6  2.00e6  2.05e6  2.08e6  2.12e6  2.16e6
## 10 Aruba   3.07e4  3.06e4  3.05e4  3.04e4  3.03e4  3.03e4  3.06e4  3.09e4  3.14e4
## # ... with 199 more rows, and 28 more variables: `1984` <dbl>,
## #   `1985` <dbl>, `1986` <dbl>, `1987` <dbl>, `1988` <dbl>, `1989` <dbl>,
## #   `1990` <dbl>, `1991` <dbl>, `1992` <dbl>, `1993` <dbl>, `1994` <dbl>,
## #   `1995` <dbl>, `1996` <dbl>, `1997` <dbl>, `1998` <dbl>, `1999` <dbl>,
## #   `2000` <dbl>, `2001` <dbl>, `2002` <dbl>, `2003` <dbl>, `2004` <dbl>,
## #   `2005` <dbl>, `2006` <dbl>, `2007` <dbl>, `2008` <dbl>, `2009` <dbl>,
## #   `2010` <dbl>, `2011` <dbl>
```

Import data using haven

1. Load the haven package.

```
library("haven")
```

2. In this exercise, you will work with data on yearly import and export numbers of sugar, both in USD and in weight. The data is given in **trade.dta**. Load the data using **read_dta** and have a look at the structure. Convert the values in Date column to dates.

```
sugar <- read_dta("data/trade.dta")
sugar
```

```
## # A tibble: 10 x 5
##   Date      Import Weight_I   Export Weight_E
##   <dbl+lbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 10      37664782 54029106 54505513 93350013
## 2 9       16316512 21584365 102700010 158000010
## 3 8       11082246 14526089 37935000 88000000
## 4 7       35677943 55034932 48515008 112000005
## 5 6       9879878 14806865 71486545 131800000
## 6 5       1539992 1749318 12311696 18500014
## 7 4        28021 54567 16489813 39599944
## 8 3         2652 3821 29273920 102072480
## 9 2       7067402 23722957 46497438 147583380
## 10 1      1033672 1964980 27131638 78268792
```

```
str(sugar) # Structure of sugar
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame': 10 obs. of 5 variables:
## $ Date : 'labelled' num 10 9 8 7 6 5 4 3 2 1
## .. attr(*, "label")= chr "Date"
## .. attr(*, "format.stata")= chr "%9.0g"
## .. attr(*, "labels")= Named num 1 2 3 4 5 6 7 8 9 10
## .. ..- attr(*, "names")= chr "2004-12-31" "2005-12-31" "2006-12-31" "2007-12-31" ...
## $ Import : num 37664782 16316512 11082246 35677943 9879878 ...
## .. attr(*, "label")= chr "Import"
## .. attr(*, "format.stata")= chr "%9.0g"
## $ Weight_I: num 54029106 21584365 14526089 55034932 14806865 ...
## .. attr(*, "label")= chr "Weight_I"
## .. attr(*, "format.stata")= chr "%9.0g"
## $ Export : num 5.45e+07 1.03e+08 3.79e+07 4.85e+07 7.15e+07 ...
## .. attr(*, "label")= chr "Export"
## .. attr(*, "format.stata")= chr "%9.0g"
## $ Weight_E: num 9.34e+07 1.58e+08 8.80e+07 1.12e+08 1.32e+08 ...
## .. attr(*, "label")= chr "Weight_E"
## .. attr(*, "format.stata")= chr "%9.0g"
## - attr(*, "label")= chr "Written by R."
```

```
sugar$Date <- as.Date(as_factor(sugar$Date))
str(sugar)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame': 10 obs. of 5 variables:
## $ Date : Date, format: "2013-12-31" "2012-12-31" ...
## $ Import : num 37664782 16316512 11082246 35677943 9879878 ...
## .. attr(*, "label")= chr "Import"
## .. attr(*, "format.stata")= chr "%9.0g"
## $ Weight_I: num 54029106 21584365 14526089 55034932 14806865 ...
## .. attr(*, "label")= chr "Weight_I"
## .. attr(*, "format.stata")= chr "%9.0g"
## $ Export : num 5.45e+07 1.03e+08 3.79e+07 4.85e+07 7.15e+07 ...
## .. attr(*, "label")= chr "Export"
## .. attr(*, "format.stata")= chr "%9.0g"
## $ Weight_E: num 9.34e+07 1.58e+08 8.80e+07 1.12e+08 1.32e+08 ...
## .. attr(*, "label")= chr "Weight_E"
## .. attr(*, "format.stata")= chr "%9.0g"
## - attr(*, "label")= chr "Written by R."
```

Import data using foreign

1. Load the foreign package.

```
library(foreign)
```

1. In this exercise, you will import data on the US presidential elections in the year 2000. The data in **florida.dta** contains the total numbers of votes for each of the four candidates as well as the total number of votes per election area in the state of Florida. Import **florida.dta** and name the resulting data frame **florida**.

```
florida <- read.dta("data/florida.dta")
tail(florida)
```

```
##      gore  bush buchanan nader  total
## 62  2647  4051      27    59   6784
## 63  1399  2326      26    29   3780
## 64 97063 82214     396   2436 182109
## 65  3835  4511      46   149   8541
## 66  5637 12176     120   265  18198
## 67  2796  4983      88    93   7960
```

2. The arguments you will use most often are `convert.dates`, `convert.factors`, `missing.type` and `convert.underscore`. Consider the dataset **edequality** which contains socio-economic measures and access to education for different individuals (source: Worldbank).

```
edu_equal_1 <- read.dta("data/edequality.dta")
str(edu_equal_1)
```

```
## 'data.frame': 12214 obs. of 27 variables:
## $ hhid : num 1 1 1 2 2 3 4 4 5 6 ...
## $ hhweight : num 627 627 627 627 627 ...
## $ location : Factor w/ 2 levels "urban location",...: 1 1 1 1 1 2 2 2 1 1 ...
## $ region : Factor w/ 9 levels "Sofia city","Bourgass",...: 8 8 8 9 9 4 4 4 8 8 ...
## $ ethnicity_head : Factor w/ 4 levels "Bulgaria","Turks",...: 2 2 2 1 1 1 1 1 1 1 ...
## $ age : num 37 11 8 73 70 75 79 80 82 83 ...
## $ gender : Factor w/ 2 levels "male","female": 2 2 1 1 2 1 1 2 2 2 ...
## $ relation : Factor w/ 9 levels "head",...: 1 3 3 1 2 1 1 2 1 1 ...
## $ literate : Factor w/ 2 levels "no","yes": 1 2 2 2 2 2 2 2 2 2 ...
## $ income_mnt : num 13.3 13.3 13.3 142.5 142.5 ...
## $ income : num 160 160 160 1710 1710 ...
## $ aggregate : num 1042 1042 1042 3271 3271 ...
## $ aggr_ind_annual : num 347 347 347 1635 1635 ...
## $ educ_completed : int 2 4 4 4 3 3 3 3 4 4 ...
## $ grade_complete : num 4 3 0 3 4 4 4 4 5 5 ...
## $ grade_all : num 4 11 8 11 8 8 8 8 13 13 ...
## $ unemployed : int 2 1 1 1 1 1 1 1 1 1 ...
## $ reason_OLF : int NA NA NA 3 3 3 9 9 3 3 ...
## $ sector : int NA NA NA NA NA NA 1 1 NA NA ...
## $ occupation : int NA NA NA NA NA NA 5 5 NA NA ...
## $ earn_mont : num 0 0 0 0 0 0 20 20 0 0 ...
## $ earn_ann : num 0 0 0 0 0 0 240 240 0 0 ...
## $ hours_week : num NA NA NA NA NA NA 30 35 NA NA ...
## $ hours_mnt : num NA NA NA NA NA ...
## $ fulltime : int NA NA NA NA NA 1 1 NA NA ...
## $ hhexp : num 100 100 100 343 343 ...
```

```

## $ legacy_pension_amt: num  NA NA NA NA NA NA NA NA NA NA ...
## - attr(*, "datalabel")= chr ""
## - attr(*, "time.stamp")= chr ""
## - attr(*, "formats")= chr  "%9.0g" "%9.0g" "%9.0g" "%9.0g" ...
## - attr(*, "types")= int   100 100 108 108 108 100 108 108 108 100 ...
## - attr(*, "val.labels")= chr  "" "" "location" "region" ...
## - attr(*, "var.labels")= chr  "hhid" "hhweight" "location" "region" ...
## - attr(*, "expansion.fields")=List of 12
## ..$ : chr  "_dta" "_svy_su1" "cluster"
## ..$ : chr  "_dta" "_svy_strata1" "strata"
## ..$ : chr  "_dta" "_svy_stages" "1"
## ..$ : chr  "_dta" "_svy_version" "2"
## ..$ : chr  "_dta" "__XijVarLabcons" "(sum) cons"
## ..$ : chr  "_dta" "ReS_Xij" "cons"
## ..$ : chr  "_dta" "ReS_str" "0"
## ..$ : chr  "_dta" "ReS_j" "group"
## ..$ : chr  "_dta" "ReS_ver" "v.2"
## ..$ : chr  "_dta" "ReS_i" "hhid dur"
## ..$ : chr  "_dta" "note1" "variables g1pc, g2pc, g3pc, g4pc, g5pc, g7pc, g8pc, g9pc, g10pc, g11pc,
## ..$ : chr  "_dta" "note0" "1"
## - attr(*, "version")= int 7
## - attr(*, "label.table")=List of 12
## ..$ location: Named int 1 2
## .. ..- attr(*, "names")= chr  "urban location" "rural location"
## ..$ region : Named int 1 2 3 4 5 6 7 8 9
## .. ..- attr(*, "names")= chr  "Sofia city" "Bourgass" "Varna" "Lovetch" ...
## ..$ ethnic : Named int 1 2 3 4
## .. ..- attr(*, "names")= chr  "Bulgaria" "Turks" "Roma" "Other"
## ..$ s2_q2 : Named int 1 2
## .. ..- attr(*, "names")= chr  "male" "female"
## ..$ s2_q3 : Named int 1 2 3 4 5 6 7 8 9
## .. ..- attr(*, "names")= chr  "head" "spouse/partner" "child
## ..$ lit : Named int 1 2
## .. ..- attr(*, "names")= chr  "no" "yes"
## ..$ : Named int 1 2 3 4
## .. ..- attr(*, "names")= chr  "never attended" "primary" "secondary" "postsecondary"
## ..$ : Named int 1 2
## .. ..- attr(*, "names")= chr  "Not unemployed" "Unemployed"
## ..$ : Named int 1 2 3 4 5 6 7 8 9 10
## .. ..- attr(*, "names")= chr  "student" "housewife/childcare" "in retirement" "illness, disability
## ..$ : Named int 1 2 3 4 5 6 7 8 9 10
## .. ..- attr(*, "names")= chr  "agriculture" "mining" "manufacturing" "utilities" ...
## ..$ : Named int 1 2 3 4 5
## .. ..- attr(*, "names")= chr  "private company" "public works program" "government,public sector,
## ..$ : Named int 1 2
## .. ..- attr(*, "names")= chr  "no" "yes"

```

```

edu_equal_2 <- read.dta("data/edequality.dta", convert.factors = FALSE)
str(edu_equal_2)

```

```

## 'data.frame': 12214 obs. of 27 variables:
## $ hhid : num 1 1 1 2 2 3 4 4 5 6 ...
## $ hhweight : num 627 627 627 627 627 ...
## $ location : int 1 1 1 1 1 2 2 2 1 1 ...
## $ region : int 8 8 8 9 9 4 4 4 8 8 ...

```

```

## $ ethnicity_head : int 2 2 2 1 1 1 1 1 1 1 ...
## $ age : num 37 11 8 73 70 75 79 80 82 83 ...
## $ gender : int 2 2 1 1 2 1 1 2 2 2 ...
## $ relation : int 1 3 3 1 2 1 1 2 1 1 ...
## $ literate : int 1 2 2 2 2 2 2 2 2 2 ...
## $ income_mnt : num 13.3 13.3 13.3 142.5 142.5 ...
## $ income : num 160 160 160 1710 1710 ...
## $ aggregate : num 1042 1042 1042 3271 3271 ...
## $ aggr_ind_annual : num 347 347 347 1635 1635 ...
## $ educ_completed : int 2 4 4 4 3 3 3 3 4 4 ...
## $ grade_complete : num 4 3 0 3 4 4 4 4 5 5 ...
## $ grade_all : num 4 11 8 11 8 8 8 8 13 13 ...
## $ unemployed : int 2 1 1 1 1 1 1 1 1 1 ...
## $ reason_OLF : int NA NA NA 3 3 3 9 9 3 3 ...
## $ sector : int NA NA NA NA NA NA 1 1 NA NA ...
## $ occupation : int NA NA NA NA NA NA 5 5 NA NA ...
## $ earn_mont : num 0 0 0 0 0 0 20 20 0 0 ...
## $ earn_ann : num 0 0 0 0 0 0 240 240 0 0 ...
## $ hours_week : num NA NA NA NA NA NA 30 35 NA NA ...
## $ hours_mnt : num NA NA NA NA NA ...
## $ fulltime : int NA NA NA NA NA NA 1 1 NA NA ...
## $ hhexp : num 100 100 100 343 343 ...
## $ legacy_pension_amt: num NA NA NA NA NA NA NA NA NA NA ...
## - attr(*, "datalabel")= chr ""
## - attr(*, "time.stamp")= chr ""
## - attr(*, "formats")= chr "%9.0g" "%9.0g" "%9.0g" "%9.0g" ...
## - attr(*, "types")= int 100 100 108 108 108 100 108 108 108 100 ...
## - attr(*, "val.labels")= chr "" "" "location" "region" ...
## - attr(*, "var.labels")= chr "hhid" "hhweight" "location" "region" ...
## - attr(*, "expansion.fields")=List of 12
## ..$ : chr "_dta" "_svy_su1" "cluster"
## ..$ : chr "_dta" "_svy_strata1" "strata"
## ..$ : chr "_dta" "_svy_stages" "1"
## ..$ : chr "_dta" "_svy_version" "2"
## ..$ : chr "_dta" "__XijVarLabcons" "(sum) cons"
## ..$ : chr "_dta" "ReS_Xij" "cons"
## ..$ : chr "_dta" "ReS_str" "0"
## ..$ : chr "_dta" "ReS_j" "group"
## ..$ : chr "_dta" "ReS_ver" "v.2"
## ..$ : chr "_dta" "ReS_i" "hhid dur"
## ..$ : chr "_dta" "note1" "variables g1pc, g2pc, g3pc, g4pc, g5pc, g7pc, g8pc, g9pc, g10pc, g11pc,
## ..$ : chr "_dta" "note0" "1"
## - attr(*, "version")= int 7
## - attr(*, "label.table")=List of 12
## ..$ location: Named int 1 2
## .. ..- attr(*, "names")= chr "urban location" "rural location"
## ..$ region : Named int 1 2 3 4 5 6 7 8 9
## .. ..- attr(*, "names")= chr "Sofia city" "Bourgass" "Varna" "Lovetch" ...
## ..$ ethnic : Named int 1 2 3 4
## .. ..- attr(*, "names")= chr "Bulgaria" "Turks" "Roma" "Other"
## ..$ s2_q2 : Named int 1 2
## .. ..- attr(*, "names")= chr "male" "female"
## ..$ s2_q3 : Named int 1 2 3 4 5 6 7 8 9
## .. ..- attr(*, "names")= chr "head" "spouse/partner" "child

```

```
## ..$ lit      : Named int 1 2
## .. ..- attr(*, "names")= chr "no" "yes"
## ..$          : Named int 1 2 3 4
## .. ..- attr(*, "names")= chr "never attended" "primary" "secondary" "postsecondary"
## ..$          : Named int 1 2
## .. ..- attr(*, "names")= chr "Not unemployed" "Unemployed"
## ..$          : Named int 1 2 3 4 5 6 7 8 9 10
## .. ..- attr(*, "names")= chr "student" "housewife/childcare" "in retirement" "illness, disability"
## ..$          : Named int 1 2 3 4 5 6 7 8 9 10
## .. ..- attr(*, "names")= chr "agriculture" "mining" "manufacturing" "utilities" ...
## ..$          : Named int 1 2 3 4 5
## .. ..- attr(*, "names")= chr "private company" "public works program" "government,public sector, ..."
## ..$          : Named int 1 2
## .. ..- attr(*, "names")= chr "no" "yes"
```

```
edu_equal_3 <- read.dta("data/edequality.dta", convert.underscore = TRUE)
str(edu_equal_3)
```

```
## 'data.frame': 12214 obs. of 27 variables:
## $ hhid : num 1 1 1 2 2 3 4 4 5 6 ...
## $ hhweight : num 627 627 627 627 627 ...
## $ location : Factor w/ 2 levels "urban location",...: 1 1 1 1 1 2 2 2 1 1 ...
## $ region : Factor w/ 9 levels "Sofia city","Bourgass",...: 8 8 8 9 9 4 4 4 8 8 ...
## $ ethnicity.head : Factor w/ 4 levels "Bulgaria","Turks",...: 2 2 2 1 1 1 1 1 1 1 ...
## $ age : num 37 11 8 73 70 75 79 80 82 83 ...
## $ gender : Factor w/ 2 levels "male","female": 2 2 1 1 2 1 1 2 2 2 ...
## $ relation : Factor w/ 9 levels "head",...: 1 3 3 1 2 1 1 2 1 1 ...
## $ literate : Factor w/ 2 levels "no","yes": 1 2 2 2 2 2 2 2 2 2 ...
## $ income.mnt : num 13.3 13.3 13.3 142.5 142.5 ...
## $ income : num 160 160 160 1710 1710 ...
## $ aggregate : num 1042 1042 1042 3271 3271 ...
## $ aggr.ind.annual : num 347 347 347 1635 1635 ...
## $ educ.completed : int 2 4 4 4 3 3 3 4 4 ...
## $ grade.complete : num 4 3 0 3 4 4 4 4 5 5 ...
## $ grade.all : num 4 11 8 11 8 8 8 8 13 13 ...
## $ unemployed : int 2 1 1 1 1 1 1 1 1 1 ...
## $ reason.OLF : int NA NA NA 3 3 3 9 9 3 3 ...
## $ sector : int NA NA NA NA NA NA 1 1 NA NA ...
## $ occupation : int NA NA NA NA NA NA 5 5 NA NA ...
## $ earn.mont : num 0 0 0 0 0 0 20 20 0 0 ...
## $ earn.ann : num 0 0 0 0 0 0 240 240 0 0 ...
## $ hours.week : num NA NA NA NA NA NA 30 35 NA NA ...
## $ hours.mnt : num NA NA NA NA NA ...
## $ fulltime : int NA NA NA NA NA NA 1 1 NA NA ...
## $ hhexp : num 100 100 100 343 343 ...
## $ legacy.pension.amt: num NA NA NA NA NA NA NA NA NA NA ...
## - attr(*, "datalabel")= chr ""
## - attr(*, "time.stamp")= chr ""
## - attr(*, "formats")= chr "%9.0g" "%9.0g" "%9.0g" "%9.0g" ...
## - attr(*, "types")= int 100 100 108 108 108 100 108 108 108 100 ...
## - attr(*, "val.labels")= chr "" "" "location" "region" ...
## - attr(*, "var.labels")= chr "hhid" "hhweight" "location" "region" ...
## - attr(*, "expansion.fields")=List of 12
## ..$ : chr "_dta" "_svy_sui" "cluster"
## ..$ : chr "_dta" "_svy_strata1" "strata"
```

```

## ..$ : chr "_dta" "_svy_stages" "1"
## ..$ : chr "_dta" "_svy_version" "2"
## ..$ : chr "_dta" "__XijVarLabcons" "(sum) cons"
## ..$ : chr "_dta" "ReS_Xij" "cons"
## ..$ : chr "_dta" "ReS_str" "0"
## ..$ : chr "_dta" "ReS_j" "group"
## ..$ : chr "_dta" "ReS_ver" "v.2"
## ..$ : chr "_dta" "ReS_i" "hhid dur"
## ..$ : chr "_dta" "note1" "variables g1pc, g2pc, g3pc, g4pc, g5pc, g7pc, g8pc, g9pc, g10pc, g11pc,
## ..$ : chr "_dta" "note0" "1"
## - attr(*, "version")= int 7
## - attr(*, "label.table")=List of 12
## ..$ location: Named int 1 2
## .. ..- attr(*, "names")= chr "urban location" "rural location"
## ..$ region : Named int 1 2 3 4 5 6 7 8 9
## .. ..- attr(*, "names")= chr "Sofia city" "Bourgass" "Varna" "Lovetch" ...
## ..$ ethnic : Named int 1 2 3 4
## .. ..- attr(*, "names")= chr "Bulgaria" "Turks" "Roma" "Other"
## ..$ s2_q2 : Named int 1 2
## .. ..- attr(*, "names")= chr "male" "female"
## ..$ s2_q3 : Named int 1 2 3 4 5 6 7 8 9
## .. ..- attr(*, "names")= chr "head" "spouse/partner" "child"
## ..$ lit : Named int 1 2
## .. ..- attr(*, "names")= chr "no" "yes"
## ..$ : Named int 1 2 3 4
## .. ..- attr(*, "names")= chr "never attanded" "primary" "secondary" "postsecondary"
## ..$ : Named int 1 2
## .. ..- attr(*, "names")= chr "Not unemployed" "Unemployed"
## ..$ : Named int 1 2 3 4 5 6 7 8 9 10
## .. ..- attr(*, "names")= chr "student" "housewife/childcare" "in retirement" "illness, disability"
## ..$ : Named int 1 2 3 4 5 6 7 8 9 10
## .. ..- attr(*, "names")= chr "agriculture" "mining" "manufacturing" "utilities" ...
## ..$ : Named int 1 2 3 4 5
## .. ..- attr(*, "names")= chr "private company" "public works program" "government,public sector, ..."
## ..$ : Named int 1 2
## .. ..- attr(*, "names")= chr "no" "yes"

```

Indexing vectors

Define the following vectors

$$x = \begin{pmatrix} 1 \\ 1.1 \\ 9 \\ 8 \\ 1 \\ 4 \\ 4 \\ 1 \end{pmatrix}, \quad y = \begin{pmatrix} 1 \\ 2 \\ 3 \\ 4 \\ 4 \\ 3 \\ 2 \\ NA \end{pmatrix}, \quad z = \begin{pmatrix} TRUE \\ TRUE \\ FALSE \\ FALSE \\ TRUE \\ FALSE \\ FALSE \\ FALSE \end{pmatrix}$$

```
x <- c(1, 1.1, 9, 8, 1, 4, 4, 1)
y <- c(1, 2, 3, 4, 4, 3, 2, NA)
z <- c(TRUE, TRUE, FALSE, FALSE, TRUE, FALSE, FALSE, FALSE)
```

1. Predict what the following commands will return (and then check if you are right):

```
x[-2]

## [1] 1 9 8 1 4 4 1

x[2:5]

## [1] 1.1 9.0 8.0 1.0

x[c(1,5,8)]

## [1] 1 1 1

x[-c(1,5,8)]

## [1] 1.1 9.0 8.0 4.0 4.0

x[y]

## [1] 1.0 1.1 9.0 8.0 8.0 9.0 1.1 NA

x[seq(2,8,by=2)]

## [1] 1.1 8.0 4.0 1.0

x[rep(1:3,4)]

## [1] 1.0 1.1 9.0 1.0 1.1 9.0 1.0 1.1 9.0 1.0 1.1 9.0
```

2. Predict what the following commands will return (and then check if you are right):

```
y[z]

## [1] 1 2 4

y[!z]

## [1] 3 4 3 2 NA

y[x>2]

## [1] 3 4 3 2

y[x==1]

## [1] 1 4 NA

y

## [1] 1 2 3 4 4 3 2 NA

x[!is.na(y)]

## [1] 1.0 1.1 9.0 8.0 1.0 4.0 4.0

y[!is.na(y)]

## [1] 1 2 3 4 4 3 2
```

3. Indexing is not only used to read certain elements of a vector but also to change them. Execute `x2 <- x` to make a copy of `x`. Change all elements of `x2` that have the value 4 to the value -4. Print `x2`.


```
x2 <- x
x2[x2 == 4] <- -4
print(x2)
```

```
## [1] 1.0 1.1 9.0 8.0 1.0 -4.0 -4.0 1.0
```

4. Change all elements of x2 that have the value 1 to a missing value (NA). Print x2.

```
x2[x2 == 1] <- NA
print(x2)
```

```
## [1] NA 1.1 9.0 8.0 NA -4.0 -4.0 NA
```

5. Execute x2[z] <- 0. Print x2.

```
x2[z] <- 0
print(x2)
```

```
## [1] 0 0 9 8 0 -4 -4 NA
```

Indexing matrices

Define the matrix x <- matrix(c(1:12,12:1),4,6).

```
x <- matrix(c(1:12, 12:1), 4, 6)
```

1. Predict what the following commands will return (and then check if you are right):

```
x[1,3]
```

```
## [1] 9
```

```
x[,5]
```

```
## [1] 8 7 6 5
```

```
x[2,]
```

```
## [1] 2 6 10 11 7 3
```

```
x[, -3]
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    5   12    8    4
## [2,]    2    6   11    7    3
## [3,]    3    7   10    6    2
## [4,]    4    8    9    5    1
```

```
x[-4,]
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]
## [1,]    1    5    9   12    8    4
## [2,]    2    6   10   11    7    3
## [3,]    3    7   11   10    6    2
```

```
x[2:3,3:4]
```

```
##      [,1] [,2]
## [1,]   10   11
```

```
## [2,] 11 10
```

```
x[2:4,4]
```

```
## [1] 11 10 9
```

2. Predict what the following commands will return (and then check if you are right):

```
x[x>5]
```

```
## [1] 6 7 8 9 10 11 12 12 11 10 9 8 7 6
```

```
x[,x[1,]<=5]
```

```
##      [,1] [,2] [,3]
```

```
## [1,] 1 5 4
```

```
## [2,] 2 6 3
```

```
## [3,] 3 7 2
```

```
## [4,] 4 8 1
```

```
x[x[,2]>6,]
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]
```

```
## [1,] 3 7 11 10 6 2
```

```
## [2,] 4 8 12 9 5 1
```

```
x[x[,2]>6,4:6]
```

```
##      [,1] [,2] [,3]
```

```
## [1,] 10 6 2
```

```
## [2,] 9 5 1
```

```
x[x[,1]<3 & x[,2]<6,]
```

```
## [1] 1 5 9 12 8 4
```

3. Print all rows where column 5 is at least three times larger than column 6.

```
x[x[, 5]>= (3 * x[, 6]) , ]
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]
```

```
## [1,] 3 7 11 10 6 2
```

```
## [2,] 4 8 12 9 5 1
```

4. Count the number of elements of `x` that are larger than 7.

```
length(x[x > 7]) # or: sum(x>7)
```

```
## [1] 10
```

5. Count the number of elements in row 2 that are smaller than their neighbors in row 1.

```
sum(x[2, ] < x[1, ])
```

```
## [1] 3
```

6. Count the number of elements of `x` that are larger than their left neighbor.

```
sum(x[, 2:6] > x[, 1:5]) # alternativ: sum(x[,-1]>x[,-6])
```

```
## [1] 10
```

Indexing dataframes

Load the data set **bsp2.txt** as data frame **bsp2** and print it.

```
bsp2 <- read.csv("data/bsp2.txt", as.is=TRUE)
print(bsp2)
```

```
##      X      Y      Z U V
## 1  2.411 2.317 5.209 B M
## 2  2.469 2.116 2.566 D M
## 3  1.066 5.471 4.856 D M
## 4  2.264 2.107 2.647 C M
## 5  2.775 3.136 3.221 C M
## 6  2.542 3.072 2.999 A M
## 7  2.200 3.272 3.174 D F
## 8  8.947 3.058 3.178 B M
## 9  5.784 2.092 3.364 B F
## 10 2.180 2.745 3.731 B M
## 11 1.580 2.008 4.463 B M
## 12 3.590 2.214 1.556 C F
## 13 4.316 1.868 0.731 D M
## 14 2.509 6.982 1.265 B M
## 15 1.436 4.650 2.002 C F
## 16 2.275 3.761 1.376 D F
## 17 3.648 1.688 1.837 C F
## 18 1.724 1.857 4.980 B M
## 19 1.785 4.424 5.611 C F
## 20 2.133 3.958 3.530 A M
## 21 5.512 1.658 3.292 B M
## 22 3.140 6.972 4.354 B M
## 23 4.610 2.740 1.497 C M
## 24 2.575 2.967 5.191 A M
## 25 1.951 1.996 1.916 A M
## 26 3.282 6.299 3.116 D M
## 27 1.022 1.764 2.853 A M
## 28 1.344 4.229 2.308 A F
## 29 1.855 4.866 3.305 A F
## 30 0.886 2.795 4.485 A F
## 31 1.541 1.507 1.996 C F
## 32 5.492 1.202 3.887 B F
## 33 0.884 3.165 3.752 D F
## 34 4.371 1.234 2.591 C M
## 35 2.638 2.195 0.804 D M
## 36 2.606 3.889 2.388 A F
## 37 0.821 1.927 3.168 D M
## 38 1.889 4.294 2.137 C M
## 39 5.683 2.467 2.356 C F
## 40 1.624 2.650 3.247 A F
```

1. Use different ways to print the second column of the data frame **bsp2** (as a vector or a data frame).

```
bsp2[, 2]
bsp2$Y
bsp2[[2]]
bsp2["Y"]
```

2. Use different ways to print columns U and V .

```
bsp2[, 4]
bsp2$U
bsp2[[4]]
bsp2["U"]
bsp2[, 5]
bsp2$V
bsp2[[5]]
bsp2["V"]
```

3. Use the `attach` command to make the variables directly accessible. Print `X`. Now `detach` the data frame again.

```
attach(bsp2)
print(X) # For safety reasons apply rm(list = ls()) upfront
detach(bsp2)
```

4. Print all rows of `bsp2` where the variable U has value A or B.

```
bsp2[bsp2$U == "A" | bsp2$U == "B", ]
```

```
##      X      Y      Z U V
## 1  2.411 2.317 5.209 B M
## 6  2.542 3.072 2.999 A M
## 8  8.947 3.058 3.178 B M
## 9  5.784 2.092 3.364 B F
## 10 2.180 2.745 3.731 B M
## 11 1.580 2.008 4.463 B M
## 14 2.509 6.982 1.265 B M
## 18 1.724 1.857 4.980 B M
## 20 2.133 3.958 3.530 A M
## 21 5.512 1.658 3.292 B M
## 22 3.140 6.972 4.354 B M
## 24 2.575 2.967 5.191 A M
## 25 1.951 1.996 1.916 A M
## 27 1.022 1.764 2.853 A M
## 28 1.344 4.229 2.308 A F
## 29 1.855 4.866 3.305 A F
## 30 0.886 2.795 4.485 A F
## 32 5.492 1.202 3.887 B F
## 36 2.606 3.889 2.388 A F
## 40 1.624 2.650 3.247 A F
```

5. Print all rows of `bsp2` where the variable X is smaller than its median and the variable Y is larger than its median.

```
bsp2[bsp2$X < median(bsp2$X) & bsp2$Y > median(bsp2$Y), ]
```

```
##      X      Y      Z U V
## 3  1.066 5.471 4.856 D M
## 7  2.200 3.272 3.174 D F
## 10 2.180 2.745 3.731 B M
## 15 1.436 4.650 2.002 C F
## 16 2.275 3.761 1.376 D F
## 19 1.785 4.424 5.611 C F
## 20 2.133 3.958 3.530 A M
```

```
## 28 1.344 4.229 2.308 A F
## 29 1.855 4.866 3.305 A F
## 30 0.886 2.795 4.485 A F
## 33 0.884 3.165 3.752 D F
## 38 1.889 4.294 2.137 C M
```

6. One can add row names to a data frame. Execute the following command and print the data frame to have a look at the new row names:

```
row.names(bsp2) <- paste(rep(LETTERS[1:20], each = 2), rep(1:2, 20), sep = "")
print(bsp2)
```

```
##      X      Y      Z U V
## A1 2.411 2.317 5.209 B M
## A2 2.469 2.116 2.566 D M
## B1 1.066 5.471 4.856 D M
## B2 2.264 2.107 2.647 C M
## C1 2.775 3.136 3.221 C M
## C2 2.542 3.072 2.999 A M
## D1 2.200 3.272 3.174 D F
## D2 8.947 3.058 3.178 B M
## E1 5.784 2.092 3.364 B F
## E2 2.180 2.745 3.731 B M
## F1 1.580 2.008 4.463 B M
## F2 3.590 2.214 1.556 C F
## G1 4.316 1.868 0.731 D M
## G2 2.509 6.982 1.265 B M
## H1 1.436 4.650 2.002 C F
## H2 2.275 3.761 1.376 D F
## I1 3.648 1.688 1.837 C F
## I2 1.724 1.857 4.980 B M
## J1 1.785 4.424 5.611 C F
## J2 2.133 3.958 3.530 A M
## K1 5.512 1.658 3.292 B M
## K2 3.140 6.972 4.354 B M
## L1 4.610 2.740 1.497 C M
## L2 2.575 2.967 5.191 A M
## M1 1.951 1.996 1.916 A M
## M2 3.282 6.299 3.116 D M
## N1 1.022 1.764 2.853 A M
## N2 1.344 4.229 2.308 A F
## O1 1.855 4.866 3.305 A F
## O2 0.886 2.795 4.485 A F
## P1 1.541 1.507 1.996 C F
## P2 5.492 1.202 3.887 B F
## Q1 0.884 3.165 3.752 D F
## Q2 4.371 1.234 2.591 C M
## R1 2.638 2.195 0.804 D M
## R2 2.606 3.889 2.388 A F
## S1 0.821 1.927 3.168 D M
## S2 1.889 4.294 2.137 C M
## T1 5.683 2.467 2.356 C F
## T2 1.624 2.650 3.247 A F
```

7. Use the row name and the variable name to print the value of variable Z at observation T1.

```
bsp2["T1", "Z"]
```

```
## [1] 2.356
```

8. Print the rows for observations G1 and G2.

```
bsp2[c("G1", "G2"), ]
```

```
##           X           Y           Z U V
## G1 4.316 1.868 0.731 D M
## G2 2.509 6.982 1.265 B M
```

Selection and transformation with dplyr

1. Load `dplyr` and `gapminder` package which will provide you with the **gapminder** dataset. How many observations and variables are in the dataset?

```
library("dplyr")
```

```
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library("gapminder")
```

```
gapminder
```

```
## # A tibble: 1,704 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 1,694 more rows
```

2. The `filter` verb extracts particular observations based on a condition. Use pipes (`%>%`) to select all information of the year 2007. For how many countries is there data available?

```
gapminder %>%
  filter(year == 2007)
```

```
## # A tibble: 142 x 6
##   country      continent  year lifeExp      pop gdpPercap
```

```
##      <fct>      <fct>      <int>      <dbl>      <int>      <dbl>
## 1 Afghanistan Asia        2007      43.8  31889923    975.
## 2 Albania      Europe      2007      76.4   3600523   5937.
## 3 Algeria      Africa      2007      72.3  33333216   6223.
## 4 Angola       Africa      2007      42.7  12420476   4797.
## 5 Argentina    Americas    2007      75.3  40301927  12779.
## 6 Australia    Oceania     2007      81.2  20434176  34435.
## 7 Austria      Europe      2007      79.8   8199783   36126.
## 8 Bahrain      Asia        2007      75.6    708573   29796.
## 9 Bangladesh   Asia        2007      64.1 150448339   1391.
## 10 Belgium     Europe      2007      79.4  10392226   33693.
## # ... with 132 more rows
```

3. Now choose all the data for the United States in the year 2007 using the `filter` command. How many observations do you get?

```
gapminder %>%
  filter(year == 2007, country == "United States")
```

```
## # A tibble: 1 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>    <int>    <dbl>
## 1 United States Americas    2007   78.2 301139947  42952.
```

4. `arrange` is a useful command for sorting data frames. First, sort the data frame by gdp per capita in ascending order. Second, sort the data set by gdp per capita in descending order for the year 2007 only.

```
gapminder %>%
  arrange(gdpPercap)
```

```
## # A tibble: 1,704 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>    <int>    <dbl>
## 1 Congo, Dem. Rep. Africa    2002   45.0 55379852    241.
## 2 Congo, Dem. Rep. Africa    2007   46.5 64606759    278.
## 3 Lesotho      Africa    1952   42.1  748747    299.
## 4 Guinea-Bissau Africa    1952   32.5  580653    300.
## 5 Congo, Dem. Rep. Africa    1997   42.6 47798986    312.
## 6 Eritrea      Africa    1952   35.9 1438760    329.
## 7 Myanmar      Asia     1952   36.3 20092996    331.
## 8 Lesotho      Africa    1957   45.0  813338    336.
## 9 Burundi     Africa    1952   39.0 2445618    339.
## 10 Eritrea     Africa    1957   38.0 1542611    344.
## # ... with 1,694 more rows
```

```
gapminder %>%
  filter(year == 2007) %>%
  arrange(desc(gdpPercap))
```

```
## # A tibble: 142 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>    <int>    <dbl>
## 1 Norway      Europe    2007   80.2  4627926  49357.
## 2 Kuwait      Asia     2007   77.6  2505559  47307.
## 3 Singapore   Asia     2007   80.0  4553009  47143.
## 4 United States Americas  2007   78.2 301139947 42952.
## 5 Ireland     Europe    2007   78.9  4109086  40676.
```

```
## 6 Hong Kong, China Asia      2007    82.2  6980412  39725.
## 7 Switzerland      Europe    2007    81.7  7554661  37506.
## 8 Netherlands      Europe    2007    79.8  16570613 36798.
## 9 Canada            Americas  2007    80.7  33390141 36319.
## 10 Iceland          Europe    2007    81.8   301931  36181.
## # ... with 132 more rows
```

5. `mutate` is useful whenever you want to change or add variables to your dataset. First, replace the variable `pop` (population) by dividing it by 1000000. Second, add a new variable `gdp` for total gross domestic product.

```
gapminder %>%
  mutate(pop = pop/1000000) #change variable
```

```
## # A tibble: 1,704 x 6
##   country      continent year lifeExp   pop gdpPercap
##   <fct>        <fct>    <int>  <dbl> <dbl>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8.43     779.
## 2 Afghanistan Asia      1957   30.3  9.24     821.
## 3 Afghanistan Asia      1962   32.0 10.3     853.
## 4 Afghanistan Asia      1967   34.0 11.5     836.
## 5 Afghanistan Asia      1972   36.1 13.1     740.
## 6 Afghanistan Asia      1977   38.4 14.9     786.
## 7 Afghanistan Asia      1982   39.9 12.9     978.
## 8 Afghanistan Asia      1987   40.8 13.9     852.
## 9 Afghanistan Asia      1992   41.7 16.3     649.
## 10 Afghanistan Asia      1997   41.8 22.2     635.
## # ... with 1,694 more rows
```

```
gapminder %>%
  mutate(gdp = gdpPercap*pop) #add new variable total gdp
```

```
## # A tibble: 1,704 x 7
##   country      continent year lifeExp   pop gdpPercap      gdp
##   <fct>        <fct>    <int>  <dbl> <int>    <dbl>    <dbl>
## 1 Afghanistan Asia      1952   28.8 8425333    779. 6567086330.
## 2 Afghanistan Asia      1957   30.3 9240934    821. 7585448670.
## 3 Afghanistan Asia      1962   32.0 10267083    853. 8758855797.
## 4 Afghanistan Asia      1967   34.0 11537966    836. 9648014150.
## 5 Afghanistan Asia      1972   36.1 13079460    740. 9678553274.
## 6 Afghanistan Asia      1977   38.4 14880372    786. 11697659231.
## 7 Afghanistan Asia      1982   39.9 12881816    978. 12598563401.
## 8 Afghanistan Asia      1987   40.8 13867957    852. 11820990309.
## 9 Afghanistan Asia      1992   41.7 16317921    649. 10595901589.
## 10 Afghanistan Asia      1997   41.8 22227415    635. 14121995875.
## # ... with 1,694 more rows
```

6. Which countries have the highest gdp in 2007?

```
gapminder %>%
  mutate(gdp = gdpPercap*pop) %>%
  filter(year == 2007) %>%
  arrange(desc(gdp))
```

```
## # A tibble: 142 x 7
##   country      continent year lifeExp   pop gdpPercap      gdp
##   <fct>        <fct>    <int>  <dbl> <int>    <dbl>    <dbl>
```



```
## 1 United States Americas 2007 78.2 301139947 42952. 1.29e13
## 2 China Asia 2007 73.0 1318683096 4959. 6.54e12
## 3 Japan Asia 2007 82.6 127467972 31656. 4.04e12
## 4 India Asia 2007 64.7 1110396331 2452. 2.72e12
## 5 Germany Europe 2007 79.4 82400996 32170. 2.65e12
## 6 United Kingdom Europe 2007 79.4 60776238 33203. 2.02e12
## 7 France Europe 2007 80.7 61083916 30470. 1.86e12
## 8 Brazil Americas 2007 72.4 190010647 9066. 1.72e12
## 9 Italy Europe 2007 80.5 58147733 28570. 1.66e12
## 10 Mexico Americas 2007 76.2 108700891 11978. 1.30e12
## # ... with 132 more rows
```

7. The basic use of the `summarize` verb is to turn many rows into one. Use it to output the mean and median of `lifeExp` as well as the total population in 2007 into a new data frame using pipes.

```
gapminder %>%
  filter(year==2007) %>%
  summarize(meanLifeExp = mean(lifeExp), medianLifeExp = median(lifeExp), totalPop = sum(as.numeric(pop)))

## # A tibble: 1 x 3
##   meanLifeExp medianLifeExp totalPop
##   <dbl>         <dbl>         <dbl>
## 1      67.0          71.9 6251013179
```

8. Now do the same, but for all years, using `group_by(year)`.

```
gapminder %>%
  group_by(year) %>%
  summarize(meanLifeExp = mean(lifeExp), medianLifeExp = median(lifeExp), totalPop = sum(as.numeric(pop)))

## # A tibble: 12 x 4
##   year meanLifeExp medianLifeExp totalPop
##   <int>         <dbl>         <dbl>         <dbl>
## 1 1952      49.1          45.1 2406957150
## 2 1957      51.5          48.4 2664404580
## 3 1962      53.6          50.9 2899782974
## 4 1967      55.7          53.8 3217478384
## 5 1972      57.6          56.5 3576977158
## 6 1977      59.6          59.7 3930045807
## 7 1982      61.5          62.4 4289436840
## 8 1987      63.2          65.8 4691477418
## 9 1992      64.2          67.7 5110710260
## 10 1997      65.0          69.4 5515204472
## 11 2002      65.7          70.8 5886977579
## 12 2007      67.0          71.9 6251013179
```

9. Again get the same statistics, but this time by continent for the year 2007.

```
gapminder %>%
  filter(year==2007) %>%
  group_by(continent) %>%
  summarize(meanLifeExp = mean(lifeExp), medianLifeExp = median(lifeExp), totalPop = sum(as.numeric(pop)))

## # A tibble: 5 x 4
##   continent meanLifeExp medianLifeExp totalPop
##   <fct>         <dbl>         <dbl>         <dbl>
## 1 Africa      54.8          52.9 929539692
## 2 Americas    73.6          72.9 898871184
```

```
## 3 Asia          70.7          72.4 3811953827
## 4 Europe        77.6          78.6 586098529
## 5 Oceania       80.7          80.7 24549947
```

10. Lastly, get the same statistics by continent and year

```
gapminder %>%
  group_by(year,continent) %>%
  summarize(meanLifeExp = mean(lifeExp), medianLifeExp = median(lifeExp), totalPop = sum(as.numeric(population)))
```

```
## # A tibble: 60 x 5
## # Groups:   year [12]
##   year continent meanLifeExp medianLifeExp totalPop
##   <int> <fct>      <dbl>         <dbl>      <dbl>
## 1  1952 Africa      39.1           38.8 237640501
## 2  1952 Americas    53.3           54.7 345152446
## 3  1952 Asia        46.3           44.9 1395357351
## 4  1952 Europe      64.4           65.9 418120846
## 5  1952 Oceania     69.3           69.3 10686006
## 6  1957 Africa      41.3           40.6 264837738
## 7  1957 Americas    56.0           56.1 386953916
## 8  1957 Asia        49.3           48.3 1562780599
## 9  1957 Europe      66.7           67.6 437890351
## 10 1957 Oceania     70.3           70.3 11941976
## # ... with 50 more rows
```

Graphics with ggplot

1. Load `dplyr`, `ggplot2` and `gapminder` package which will provide you with the `gapminder` dataset. How many observations and variables are in the dataset?


```
library("dplyr")
library("ggplot2")
library("gapminder")
```

2. Save all data from 2007 into the data frame `gapminder_2007`.


```
gapminder_2007 <- gapminder %>%
  filter(year == 2007)
```

3. Visualize countries wealth (`gdpPercap` on x axis) against life expectancy (`lifeExp` on y axis) using the good old `plot` command. Compare this to the way `ggplot` draws a scatterplot when using `geom_point`.

```
plot(gapminder_2007$gdpPercap, gapminder_2007$lifeExp)
```



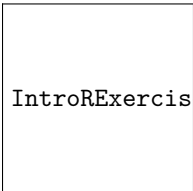
```
ggplot(gapminder_2007, aes(x = gdpPercap, y = lifeExp)) +
  geom_point()
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-108-2.pdf

4. Now add a log scale `scale_x_log10` for `gdpPercap` to your ggplot scatterplot.

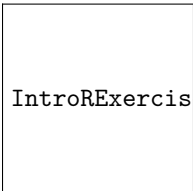
```
ggplot(gapminder_2007, aes(x = gdpPercap, y = lifeExp)) +  
  geom_point() +  
  scale_x_log10() ## each unit on the x-axis represents a change of 10 times the gdp
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-109-1.pdf

5. Create a scatter plot comparing `pop` and `gdpPercap` for the year 2007, with both axes on a log scale.

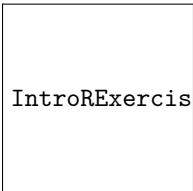
```
ggplot(gapminder_2007, aes(x = pop, y = gdpPercap)) +  
  geom_point() +  
  scale_x_log10() +  
  scale_y_log10()
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-110-1.pdf

6. To add more variables to a 2-dimensional plot, we can use two more aesthetics: color for a categorical variable and size for numerical variables. Add the continent and pop to the scatterplot of `gdpPercap` and `lifeExp` for the year 2007

```
ggplot(gapminder_2007, aes(x = gdpPercap, y = lifeExp, color = continent, size = pop)) +  
  geom_point() +  
  scale_x_log10()
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-111-1.pdf

7. Now we want to compare the dynamic relationship between `gdpPercap` and `lifeExp` for all years. Use `facet_wrap(~ year)` on the original dataset `gapminder` to add a facet to your scatterplot.

```
ggplot(gapminder, aes(x = gdpPercap, y = lifeExp, color = continent, size = pop)) +  
  geom_point() +  
  scale_x_log10() +  
  facet_wrap(~ year)
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-112-1.pdf

8. Create a data frame from the **gapminder** dataset with summarized data with the mean of `lifeExp` and total population, both grouped by year. Visualize this summarized data using `ggplot` and let your y axis begin at 0.

```
by_year <- gapminder %>%  
  group_by(year) %>%  
  summarize(meanLifeExp = mean(lifeExp), totalPop = sum(as.numeric(pop)))  
ggplot(by_year, aes(x = year, y = totalPop)) +  
  geom_point() +  
  expand_limits(y=0) # start y axis at 0
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-113-1.pdf

9. Create a data frame from the **gapminder** dataset with summarized data with the mean of `lifeExp` and total population, both grouped by year and continent. Visualize this summarized data using `ggplot` and let your y axis begin at 0.

```
by_year_continent <- gapminder %>%  
  group_by(year, continent) %>%  
  summarize(meanLifeExp = mean(lifeExp), totalPop = sum(as.numeric(pop)))  
ggplot(by_year_continent, aes(x = year, y = meanLifeExp, color = continent)) +  
  geom_point() +  
  expand_limits(y=0) # start y axis at 0
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-114-1.pdf

10. Create a line plot (`geom_line`) with `ggplot` for the just created subset of data grouped by year and continent. Put year on the x axis, `meanLifeExp` on the y axis and let the color indicate the continents.

```
ggplot(by_year_continent, aes(x = year, y = meanLifeExp, color = continent)) +  
  geom_line() +  
  expand_limits(y=0) # start y axis at 0
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-115-1.pdf

11. Filter the original **gapminder** data for the year 2007, group by continent and summarize the mean life

expectancy for this data frame. Create a bar plot (`geom_col`) with `ggplot`.

```
by_continent <- gapminder %>%  
  filter(year == 2007) %>%  
  group_by(continent) %>%  
  summarize(meanLifeExp = mean(lifeExp))  
#x is categorical variable  
ggplot(by_continent, aes(x = continent, y = meanLifeExp, color = continent)) +  
  geom_col()
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-116-1.pdf

12. Create a histogram (`geom_histogram`) of population (`pop`) in 2007 using a log scale.

```
gapminder_2007 <- gapminder %>%  
  filter(year == 2007)  
  
ggplot(gapminder_2007, aes(x = pop)) +  
  geom_histogram() +  
  scale_x_log10()
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-117-1.pdf

13. Create a boxplot (`geom_bloplot`) comparing `gdpPercap` among continents for the year 2007 and add a title to the graph using `ggtitle`.

```
gapminder_2007 <- gapminder %>%  
  filter(year == 2007)  
ggplot(gapminder_2007, aes(x = continent, y = gdpPercap)) +  
  geom_bloplot() +  
  scale_y_log10() +  
  ggtitle("Comparing GDP per capita across continents (log-scale)")
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-118-1.pdf

Histograms


In this section, please always use the command `truehist` (which is included in the `MASS` package) to generate histograms.

```
library(MASS)
```

```
##  
## Attaching package: 'MASS'  
## The following object is masked from 'package:dplyr':  
##  
## select
```


1. Load the file **gemeinden2006.csv** into a data frame. Delete all observations where the number of inhabitants (`Einwohner`) is smaller than 5. Plot the histogram of the logarithm of the variable `Einwohner`.

```
gemeinden2006 <- read.csv2("data/gemeinden2006.csv")  
gemeinden2006new <- gemeinden2006[gemeinden2006$Einwohner >= 5, ] #or x[!x$Einwohner < 5, ]  
truehist(log(gemeinden2006new$Einwohner), col = "lightblue")
```




2. Add the density function of a fitted normal distribution to the histogram.

```
truehist(log(gemeinden2006new$Einwohner), col = "lightblue")  
m <- mean(log(gemeinden2006new$Einwohner))  
s <- sd(log(gemeinden2006new$Einwohner))  
x <- seq(1, 15, length = 500)  
lines(x, dnorm(x, mean = m, sd = s), lwd = 2)
```



3. Load the Stata file **mikrozensus2002cf.dta** into a data frame. Consider the variable `ef462` (rent in April 2002). Drop all observations where the rent exceeds 2000 Euro. Plot the histogram.

```
library(foreign)  
mikrozensus2002cf <- read.dta("data/mikrozensus2002cf.dta")  
y <- mikrozensus2002cf$ef462[!is.na(mikrozensus2002cf$ef462)]  
yy <- y[y <= 2000]  
truehist(yy, col = "pink", xlab = "rent", ylab = "density", main = "histogram of rents")
```

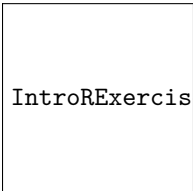


IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-123-1.pdf

4. Load the Stata file **mikrozensus2002cf.dta** into a data frame.

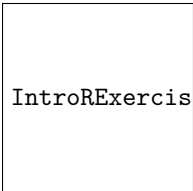
- Plot the histogram of the variable **ef453** (size of flat in square meters).
- Drop all observations with more than 300 m^2 and plot the histogram again.
- Set the number of bins in the histogram to 15.

```
truehist(mikrozensus2002cf$ef453, col = "lightblue", xlab = "size of flat (in qm)", ylab = "density")
```



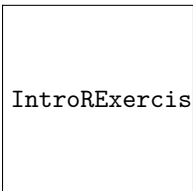
IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-124-1.pdf

```
truehist(mikrozensus2002cf$ef453[mikrozensus2002cf$ef453 <= 300], col = "lightgreen", xlab = "size of f
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-124-2.pdf

```
truehist(mikrozensus2002cf$ef453[mikrozensus2002cf$ef453 <= 300], col = "steelblue", nbins = 15, xlab =
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-124-3.pdf

Correlation and covariance

1. Execute `data(Titanic)` to load the object `Titanic` of class `table`. Print it as an ordinary table and as a flat table. Plot it as well. Compute the univariate marginal distributions using the `apply` command. Compute the bivariate marginal distribution of survival and social class (again using `apply`).

```
data(Titanic)
Titanic
```

```
## , , Age = Child, Survived = No
##
##      Sex
## Class  Male Female
```

```
## 1st      0      0
## 2nd      0      0
## 3rd     35     17
## Crew     0      0
##
## , , Age = Adult, Survived = No
##
##      Sex
## Class  Male Female
## 1st    118      4
## 2nd    154     13
## 3rd    387     89
## Crew   670      3
##
## , , Age = Child, Survived = Yes
##
##      Sex
## Class  Male Female
## 1st      5      1
## 2nd     11     13
## 3rd     13     14
## Crew     0      0
##
## , , Age = Adult, Survived = Yes
##
##      Sex
## Class  Male Female
## 1st     57    140
## 2nd     14     80
## 3rd     75     76
## Crew    192     20
```

```
table(Titanic)
```

```
## Titanic
##  0  1  3  4  5 11 13 14 17 20 35 57 75 76 80 89 118 140
##  8  1  1  1  1  1  3  2  1  1  1  1  1  1  1  1  1  1
## 154 192 387 670
##  1  1  1  1
```

```
ftable(Titanic)
```

```
##
##      Survived  No Yes
## Class Sex  Age
## 1st  Male  Child      0  5
##      Adult      118 57
##      Female Child      0  1
##      Adult      4 140
## 2nd  Male  Child      0 11
##      Adult     154 14
##      Female Child      0 13
##      Adult      13 80
## 3rd  Male  Child     35 13
##      Adult    387 75
##      Female Child     17 14
##      Adult     89 76
```

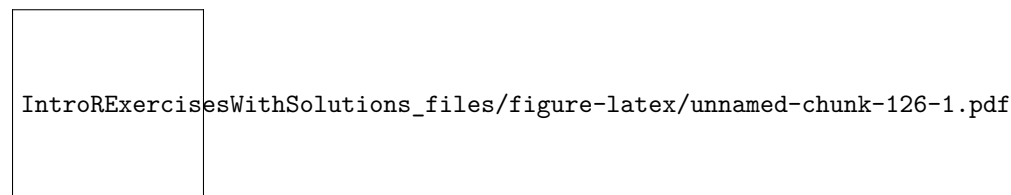


```
## Crew Male Child      0  0
##           Adult    670 192
##           Female Child      0  0
##           Adult      3  20
```

```
fable(Titanic, row.vars = c("Survived", "Age")) # Write 'Survived' and 'Age' into rows
```

```
##           Class 1st      2nd      3rd      Crew
##           Sex  Male Female Male Female Male Female Male Female
## Survived Age
## No      Child      0      0      0      0  35      17      0      0
##           Adult    118      4  154      13 387      89 670      3
## Yes     Child      5      1  11      13 13      14      0      0
##           Adult    57     140  14      80 75      76 192     20
```

```
plot(Titanic)
```



```
apply(Titanic, 1, sum)
```

```
## 1st 2nd 3rd Crew
## 325 285 706 885
```

```
apply(Titanic, 2, sum)
```

```
## Male Female
## 1731 470
```

```
apply(Titanic, 3, sum)
```

```
## Child Adult
## 109 2092
```

```
apply(Titanic, 4, sum)
```

```
## No Yes
## 1490 711
```

```
apply(Titanic, c(1, 4), sum)
```

```
##           Survived
## Class No Yes
## 1st 122 203
## 2nd 167 118
## 3rd 528 178
## Crew 673 212
```

```
# or margin.table(Titanic, c(1, 4))
```

2. Load the file `covmat.csv` into a data frame.

```
covmat <- read.csv("data/covmat.csv")
```

- Compute the covariance matrix using the option `use="complete"`. Check if the covariance matrix is positive definite.

```
cov(covmat)
```

```
##           V1 V2 V3
## V1 1.170296 NA NA
## V2      NA NA NA
## V3      NA NA NA
```

```
cov(covmat,use="complete")
```

```
##           V1           V2           V3
## V1 0.35373543 -0.09918577 -0.1629030
## V2 -0.09918577 0.34900069 0.4293324
## V3 -0.16290303 0.42933237 0.7174221
```

```
if (sum(eigen(cov(covmat,use="complete"))$value>0) == dim(covmat)[2]){
  print("Matrix is positive definite")
} else {
  print("Matrix is not positive definite")
}
```

```
## [1] "Matrix is positive definite"
```

- Now compute the covariance using the option `pairwise` and check again, if the covariance matrix is positive definite.

```
cov(covmat,use="pairwise")
```

```
##           V1           V2           V3
## V1 1.1702962 -0.1119788 0.1337978
## V2 -0.1119788 0.3051700 0.4293324
## V3 0.1337978 0.4293324 0.5038195
```

```
if (sum(eigen(cov(covmat,use="pairwise"))$value>0) == dim(covmat)[2]){
  print("Matrix is positive definite")
} else {
  print("Matrix is not positive definite")
}
```

```
## [1] "Matrix is not positive definite"
```

Cleaning data

We will consider historical weather data for Boston, USA for 12 months beginning in December 2014.

1. Load the packages `tidyr`, `dplyr`, `lubridate`, and `stringr`. Import the data using `weather <- readRDS("weather.rds")` and have a look how **dirty** it is.

```
library("tidyr"); library("dplyr"); library("lubridate"); library("stringr")
```

```
##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##      date
```

```
detach("package:MASS")
weather <- readRDS("data/weather.rds")
```

2. The first step is to understand the structure of your data with `class`, `dim`, `names`, `str`, `glimpse`, and `summary`. Also preview the first and last 15 observations with `head` and `tail`.

```
class(weather) # Verify that weather is a data.frame
```

```
## [1] "data.frame"
```

```
dim(weather) # Check the dimensions
```

```
## [1] 286 35
```

```
names(weather) # View the column names
```

```
## [1] "X"      "year"   "month"  "measure" "X1"      "X2"      "X3"
## [8] "X4"      "X5"      "X6"      "X7"      "X8"      "X9"      "X10"
## [15] "X11"     "X12"     "X13"     "X14"     "X15"     "X16"     "X17"
## [22] "X18"     "X19"     "X20"     "X21"     "X22"     "X23"     "X24"
## [29] "X25"     "X26"     "X27"     "X28"     "X29"     "X30"     "X31"
```

```
str(weather) # View the structure of the data
```

```
## 'data.frame': 286 obs. of 35 variables:
## $ X : int 1 2 3 4 5 6 7 8 9 10 ...
## $ year : int 2014 2014 2014 2014 2014 2014 2014 2014 2014 2014 ...
## $ month : int 12 12 12 12 12 12 12 12 12 12 ...
## $ measure: chr "Max.TemperatureF" "Mean.TemperatureF" "Min.TemperatureF" "Max.Dew.PointF" ...
## $ X1 : chr "64" "52" "39" "46" ...
## $ X2 : chr "42" "38" "33" "40" ...
## $ X3 : chr "51" "44" "37" "49" ...
## $ X4 : chr "43" "37" "30" "24" ...
## $ X5 : chr "42" "34" "26" "37" ...
## $ X6 : chr "45" "42" "38" "45" ...
## $ X7 : chr "38" "30" "21" "36" ...
## $ X8 : chr "29" "24" "18" "28" ...
## $ X9 : chr "49" "39" "29" "49" ...
## $ X10 : chr "48" "43" "38" "45" ...
## $ X11 : chr "39" "36" "32" "37" ...
## $ X12 : chr "39" "35" "31" "28" ...
## $ X13 : chr "42" "37" "32" "28" ...
## $ X14 : chr "45" "39" "33" "29" ...
## $ X15 : chr "42" "37" "32" "33" ...
## $ X16 : chr "44" "40" "35" "42" ...
## $ X17 : chr "49" "45" "41" "46" ...
## $ X18 : chr "44" "40" "36" "34" ...
## $ X19 : chr "37" "33" "29" "25" ...
## $ X20 : chr "36" "32" "27" "30" ...
## $ X21 : chr "36" "33" "30" "30" ...
## $ X22 : chr "44" "39" "33" "39" ...
## $ X23 : chr "47" "45" "42" "45" ...
## $ X24 : chr "46" "44" "41" "46" ...
## $ X25 : chr "59" "52" "44" "58" ...
## $ X26 : chr "50" "44" "37" "31" ...
## $ X27 : chr "52" "45" "38" "34" ...
## $ X28 : chr "52" "46" "40" "42" ...
```

```
## $ X29 : chr "41" "36" "30" "26" ...
## $ X30 : chr "30" "26" "22" "10" ...
## $ X31 : chr "30" "25" "20" "8" ...
```

```
glimpse(weather) # Look at the structure using dplyr's glimpse()
```

```
## Observations: 286
## Variables: 35
## $ X          <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, ...
## $ year       <int> 2014, 2014, 2014, 2014, 2014, 2014, 2014, 2014, 2014, 2014, ...
## $ month      <int> 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, ...
## $ measure    <chr> "Max.TemperatureF", "Mean.TemperatureF", "Min.Temperat...
## $ X1         <chr> "64", "52", "39", "46", "40", "26", "74", "63", "52", ...
## $ X2         <chr> "42", "38", "33", "40", "27", "17", "92", "72", "51", ...
## $ X3         <chr> "51", "44", "37", "49", "42", "24", "100", "79", "57", ...
## $ X4         <chr> "43", "37", "30", "24", "21", "13", "69", "54", "39", ...
## $ X5         <chr> "42", "34", "26", "37", "25", "12", "85", "66", "47", ...
## $ X6         <chr> "45", "42", "38", "45", "40", "36", "100", "93", "85", ...
## $ X7         <chr> "38", "30", "21", "36", "20", "-3", "92", "61", "29", ...
## $ X8         <chr> "29", "24", "18", "28", "16", "3", "92", "70", "47", "...
## $ X9         <chr> "49", "39", "29", "49", "41", "28", "100", "93", "86", ...
## $ X10        <chr> "48", "43", "38", "45", "39", "37", "100", "95", "89", ...
## $ X11        <chr> "39", "36", "32", "37", "31", "27", "92", "87", "82", ...
## $ X12        <chr> "39", "35", "31", "28", "27", "25", "85", "75", "64", ...
## $ X13        <chr> "42", "37", "32", "28", "26", "24", "75", "65", "55", ...
## $ X14        <chr> "45", "39", "33", "29", "27", "25", "82", "68", "53", ...
## $ X15        <chr> "42", "37", "32", "33", "29", "27", "89", "75", "60", ...
## $ X16        <chr> "44", "40", "35", "42", "36", "30", "96", "85", "73", ...
## $ X17        <chr> "49", "45", "41", "46", "41", "32", "100", "85", "70", ...
## $ X18        <chr> "44", "40", "36", "34", "30", "26", "89", "73", "57", ...
## $ X19        <chr> "37", "33", "29", "25", "22", "20", "69", "63", "56", ...
## $ X20        <chr> "36", "32", "27", "30", "24", "20", "89", "79", "69", ...
## $ X21        <chr> "36", "33", "30", "30", "27", "25", "85", "77", "69", ...
## $ X22        <chr> "44", "39", "33", "39", "34", "25", "89", "79", "69", ...
## $ X23        <chr> "47", "45", "42", "45", "42", "37", "100", "91", "82", ...
## $ X24        <chr> "46", "44", "41", "46", "44", "41", "100", "98", "96", ...
## $ X25        <chr> "59", "52", "44", "58", "43", "29", "100", "75", "49", ...
## $ X26        <chr> "50", "44", "37", "31", "29", "28", "70", "60", "49", ...
## $ X27        <chr> "52", "45", "38", "34", "31", "29", "70", "60", "50", ...
## $ X28        <chr> "52", "46", "40", "42", "35", "27", "76", "65", "53", ...
## $ X29        <chr> "41", "36", "30", "26", "20", "10", "64", "51", "37", ...
## $ X30        <chr> "30", "26", "22", "10", "4", "-6", "50", "38", "26", "...
## $ X31        <chr> "30", "25", "20", "8", "5", "1", "57", "44", "31", "30..."
```

```
summary(weather) # View a summary of the data
```

```
##           X           year           month           measure
## Min.      : 1.00      Min.    :2014      Min.    : 1.000      Length:286
## 1st Qu.: 72.25      1st Qu.:2015      1st Qu.: 4.000      Class :character
## Median :143.50      Median :2015      Median : 7.000      Mode  :character
## Mean    :143.50      Mean    :2015      Mean    : 6.923
## 3rd Qu.:214.75      3rd Qu.:2015      3rd Qu.:10.000
## Max.    :286.00      Max.    :2015      Max.    :12.000
##           X1           X2           X3
## Length:286      Length:286      Length:286
```

## Class :character	Class :character	Class :character
## Mode :character	Mode :character	Mode :character
##		
##		
##		
## X4	X5	X6
## Length:286	Length:286	Length:286
## Class :character	Class :character	Class :character
## Mode :character	Mode :character	Mode :character
##		
##		
##		
## X7	X8	X9
## Length:286	Length:286	Length:286
## Class :character	Class :character	Class :character
## Mode :character	Mode :character	Mode :character
##		
##		
##		
## X10	X11	X12
## Length:286	Length:286	Length:286
## Class :character	Class :character	Class :character
## Mode :character	Mode :character	Mode :character
##		
##		
##		
## X13	X14	X15
## Length:286	Length:286	Length:286
## Class :character	Class :character	Class :character
## Mode :character	Mode :character	Mode :character
##		
##		
##		
## X16	X17	X18
## Length:286	Length:286	Length:286
## Class :character	Class :character	Class :character
## Mode :character	Mode :character	Mode :character
##		
##		
##		
## X19	X20	X21
## Length:286	Length:286	Length:286
## Class :character	Class :character	Class :character
## Mode :character	Mode :character	Mode :character
##		
##		
##		
## X22	X23	X24
## Length:286	Length:286	Length:286
## Class :character	Class :character	Class :character
## Mode :character	Mode :character	Mode :character
##		
##		
##		

```
##      X25      X26      X27
## Length:286 Length:286 Length:286
## Class :character Class :character Class :character
## Mode :character Mode :character Mode :character
##
##
##
##      X28      X29      X30
## Length:286 Length:286 Length:286
## Class :character Class :character Class :character
## Mode :character Mode :character Mode :character
##
##
##
##      X31
## Length:286
## Class :character
## Mode :character
##
##
##
```

```
head(weather,15)
```

```
##      X year month      measure      X1      X2      X3      X4      X5
## 1  1 2014    12      Max.TemperatureF  64    42    51    43    42
## 2  2 2014    12      Mean.TemperatureF  52    38    44    37    34
## 3  3 2014    12      Min.TemperatureF  39    33    37    30    26
## 4  4 2014    12      Max.Dew.PointF    46    40    49    24    37
## 5  5 2014    12      MeanDew.PointF    40    27    42    21    25
## 6  6 2014    12      Min.DewpointF    26    17    24    13    12
## 7  7 2014    12      Max.Humidity     74    92   100    69    85
## 8  8 2014    12      Mean.Humidity     63    72    79    54    66
## 9  9 2014    12      Min.Humidity     52    51    57    39    47
## 10 10 2014    12      Max.Sea.Level.PressureIn 30.45 30.71 30.4 30.56 30.68
## 11 11 2014    12      Mean.Sea.Level.PressureIn 30.13 30.59 30.07 30.33 30.59
## 12 12 2014    12      Min.Sea.Level.PressureIn 30.01 30.4 29.87 30.09 30.45
## 13 13 2014    12      Max.VisibilityMiles    10    10    10    10    10
## 14 14 2014    12      Mean.VisibilityMiles    10     8     5    10    10
## 15 15 2014    12      Min.VisibilityMiles    10     2     1    10     5
##      X6      X7      X8      X9      X10      X11      X12      X13      X14      X15      X16      X17
## 1      45      38      29      49      48      39      39      42      45      42      44      49
## 2      42      30      24      39      43      36      35      37      39      37      40      45
## 3      38      21      18      29      38      32      31      32      33      32      35      41
## 4      45      36      28      49      45      37      28      28      29      33      42      46
## 5      40      20      16      41      39      31      27      26      27      29      36      41
## 6      36      -3       3      28      37      27      25      24      25      27      30      32
## 7     100      92      92     100     100      92      85      75      82      89      96     100
## 8      93      61      70      93      95      87      75      65      68      75      85      85
## 9      85      29      47      86      89      82      64      55      53      60      73      70
## 10 30.42 30.69 30.77 30.51 29.58 29.81 29.88 29.86 29.91 30.15 30.17 29.91
## 11 30.24 30.46 30.67 30.04 29.5 29.61 29.85 29.82 29.83 30.05 30.09 29.75
## 12 30.16 30.24 30.51 29.49 29.43 29.44 29.81 29.78 29.78 29.91 29.92 29.69
## 13    10     10     10     10     10     10     10     10     10     10     10     10
## 14     4     10      8      2       3       7     10     10     10     10      9      6
```

```
## 15      0      5      2      1      1      1      7      10      10      10      5      1
##      X18     X19     X20     X21     X22     X23     X24     X25     X26     X27     X28     X29
## 1      44      37      36      36      44      47      46      59      50      52      52      41
## 2      40      33      32      33      39      45      44      52      44      45      46      36
## 3      36      29      27      30      33      42      41      44      37      38      40      30
## 4      34      25      30      30      39      45      46      58      31      34      42      26
## 5      30      22      24      27      34      42      44      43      29      31      35      20
## 6      26      20      20      25      25      37      41      29      28      29      27      10
## 7      89      69      89      85      89      100     100     100     70      70      76      64
## 8      73      63      79      77      79      91      98      75      60      60      65      51
## 9      57      56      69      69      69      82      96      49      49      50      53      37
## 10 29.87 30.15 30.31 30.37 30.4 30.31 30.13 29.96 30.16 30.22 29.99 30.22
## 11 29.78 29.98 30.26 30.32 30.35 30.23 29.9 29.63 30.11 30.14 29.87 30.12
## 12 29.71 29.86 30.17 30.28 30.3 30.16 29.55 29.47 29.99 30.03 29.77 30
## 13      10      10      10      10      10      10      2      10      10      10      10      10
## 14      10      10      10      9      10      5      1      8      10      10      10      10
## 15      10      10      7      6      4      1      0      1      10      10      10      10
##      X30     X31
## 1      30      30
## 2      26      25
## 3      22      20
## 4      10      8
## 5      4      5
## 6      -6      1
## 7      50      57
## 8      38      44
## 9      26      31
## 10 30.36 30.32
## 11 30.32 30.25
## 12 30.23 30.13
## 13      10      10
## 14      10      10
## 15      10      10
```

```
tail(weather,15)
```

```
##      X year month      measure      X1      X2      X3      X4      X5
## 272 272 2015      12      Mean.Humidity      83 <NA> <NA> <NA> <NA>
## 273 273 2015      12      Min.Humidity      69 <NA> <NA> <NA> <NA>
## 274 274 2015      12 Max.Sea.Level.PressureIn 30.4 <NA> <NA> <NA> <NA>
## 275 275 2015      12 Mean.Sea.Level.PressureIn 30.24 <NA> <NA> <NA> <NA>
## 276 276 2015      12 Min.Sea.Level.PressureIn 30.01 <NA> <NA> <NA> <NA>
## 277 277 2015      12      Max.VisibilityMiles      10 <NA> <NA> <NA> <NA>
## 278 278 2015      12      Mean.VisibilityMiles      8 <NA> <NA> <NA> <NA>
## 279 279 2015      12      Min.VisibilityMiles      1 <NA> <NA> <NA> <NA>
## 280 280 2015      12      Max.Wind.SpeedMPH      15 <NA> <NA> <NA> <NA>
## 281 281 2015      12      Mean.Wind.SpeedMPH      6 <NA> <NA> <NA> <NA>
## 282 282 2015      12      Max.Gust.SpeedMPH      17 <NA> <NA> <NA> <NA>
## 283 283 2015      12      PrecipitationIn      0.14 <NA> <NA> <NA> <NA>
## 284 284 2015      12      CloudCover      7 <NA> <NA> <NA> <NA>
## 285 285 2015      12      Events      Rain <NA> <NA> <NA> <NA>
## 286 286 2015      12      WindDirDegrees      109 <NA> <NA> <NA> <NA>
##      X6      X7      X8      X9     X10     X11     X12     X13     X14     X15     X16     X17     X18     X19
## 272 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 273 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
```

```
## 274 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 275 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 276 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 277 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 278 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 279 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 280 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 281 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 282 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 283 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 284 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 285 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 286 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
##      X20  X21  X22  X23  X24  X25  X26  X27  X28  X29  X30  X31
## 272 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 273 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 274 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 275 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 276 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 277 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 278 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 279 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 280 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 281 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 282 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 283 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 284 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 285 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 286 <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
```

3. The **weather** dataset suffers from one of the five most common symptoms of messy data: column names are values. In particular, the column names X1-X31 represent days of the month, which should really be values of a new variable called **day**. The **tidyr** package provides the **gather()** function for exactly this scenario. Notice that **gather()** allows you to select multiple columns to be gathered by using the **:** operator. Call **gather()** on the weather data to gather columns X1-X31. The two columns created as a result should be called **day** and **value**. Save the result as **weather2** and view it with **head()**.

```
weather2 <- gather(weather, day, value, X1:X31, na.rm = TRUE)
head(weather2)
```

```
##   X year month      measure day value
## 1 1 2014    12 Max.TemperatureF X1    64
## 2 2 2014    12 Mean.TemperatureF X1    52
## 3 3 2014    12 Min.TemperatureF  X1    39
## 4 4 2014    12   Max.Dew.PointF  X1    46
## 5 5 2014    12   MeanDew.PointF  X1    40
## 6 6 2014    12   Min.DewpointF   X1    26
```

4. Our data suffer from a second common symptom of messy data: values are variable names. Specifically, values in the **measure** column should be variables (i.e. column names) in our dataset. The **spread()** function from **tidyr** is designed to help with this. Remove the first column of **weather2**, assigning to **without_x**. Spread the **measure** column of **without_x** and save the result to **weather3**. View the result with **head()**.

```
without_x <- weather2[, -1] # First remove column of row names
weather3 <- spread(without_x, measure, value) # Spread the data
```



```
head(weather3)
```

```
##   year month day CloudCover   Events Max.Dew.PointF Max.Gust.SpeedMPH
## 1 2014    12  X1         6     Rain           46             29
## 2 2014    12 X10         8     Rain           45             29
## 3 2014    12 X11         8 Rain-Snow          37             28
## 4 2014    12 X12         7     Snow           28             21
## 5 2014    12 X13         5           28             23
## 6 2014    12 X14         4           29             20
##   Max.Humidity Max.Sea.Level.PressureIn Max.TemperatureF
## 1           74              30.45              64
## 2          100              29.58              48
## 3           92              29.81              39
## 4           85              29.88              39
## 5           75              29.86              42
## 6           82              29.91              45
##   Max.VisibilityMiles Max.Wind.SpeedMPH Mean.Humidity
## 1              10              22              63
## 2              10              23              95
## 3              10              21              87
## 4              10              16              75
## 5              10              17              65
## 6              10              15              68
##   Mean.Sea.Level.PressureIn Mean.TemperatureF Mean.VisibilityMiles
## 1              30.13              52              10
## 2              29.5              43              3
## 3              29.61              36              7
## 4              29.85              35             10
## 5              29.82              37             10
## 6              29.83              39             10
##   Mean.Wind.SpeedMPH MeanDew.PointF Min.DewpointF Min.Humidity
## 1              13              40              26              52
## 2              13              39              37              89
## 3              13              31              27              82
## 4              11              27              25              64
## 5              12              26              24              55
## 6              10              27              25              53
##   Min.Sea.Level.PressureIn Min.TemperatureF Min.VisibilityMiles
## 1              30.01              39              10
## 2              29.43              38              1
## 3              29.44              32              1
## 4              29.81              31              7
## 5              29.78              32             10
## 6              29.78              33             10
##   PrecipitationIn WindDirDegrees
## 1              0.01             268
## 2              0.28             357
## 3              0.02             230
## 4              T             286
## 5              T             298
## 6              0.00             306
```

5. A good package and function to tidy up dates into the same format is `lubridate`, e.g. try out this code
#Dates with lubridate for most common combinations

```
ymd("2015-08-25")
```

```
## [1] "2015-08-25"
```

```
ymd("2015 August 25")
```

```
## [1] "2015-08-25"
```

```
mdy("August 25, 2015")
```

```
## [1] "2015-08-25"
```

```
hms("13:33:09")
```

```
## [1] "13H 33M 9S"
```

```
ymd_hms("2015/08/25 13.33.09")
```

```
## [1] "2015-08-25 13:33:09 UTC"
```

We'll start by combining the year, month, and day columns and recoding the resulting character column as a date. We can use a combination of `stringr`, and `lubridate` to accomplish this task.

- Use `stringr`'s `str_replace()` to remove the Xs from the day column of `weather3`.
- Create a new column called `date`. Use the `unite()` function from `tidyr` to paste together the year, month, and day columns in order, using `-` as a separator.
- Coerce the `date` column using the appropriate function from `lubridate`.
- Use `select()` to reorder columns, saving the result to `weather5`.
- View the head of `weather5`.

```
weather3$day <- str_replace(weather3$day, "X", "") # Remove X's from day column
weather4 <- unite(weather3, date, year, month, day, sep = "-") # Unite the year, month, and day columns
weather4$date <- ymd(weather4$date) # Convert date column to proper date format using lubridates's ymd()
weather5 <- select(weather4, date, Events, CloudCover:WindDirDegrees) # Rearrange columns using dplyr's
head(weather5) # View the head of weather5
```

```
##           date      Events CloudCover Max.Dew.PointF Max.Gust.SpeedMPH
## 1 2014-12-01      Rain          6          46          29
## 2 2014-12-10      Rain          8          45          29
## 3 2014-12-11 Rain-Snow          8          37          28
## 4 2014-12-12      Snow          7          28          21
## 5 2014-12-13          5          28          23
## 6 2014-12-14          4          29          20
## Max.Humidity Max.Sea.Level.PressureIn Max.TemperatureF
## 1          74          30.45          64
## 2          100          29.58          48
## 3          92          29.81          39
## 4          85          29.88          39
## 5          75          29.86          42
## 6          82          29.91          45
## Max.VisibilityMiles Max.Wind.SpeedMPH Mean.Humidity
## 1          10          22          63
## 2          10          23          95
## 3          10          21          87
## 4          10          16          75
## 5          10          17          65
```

```
## 6          10          15          68
## Mean.Sea.Level.PressureIn Mean.TemperatureF Mean.VisibilityMiles
## 1          30.13          52          10
## 2          29.5          43           3
## 3          29.61          36           7
## 4          29.85          35          10
## 5          29.82          37          10
## 6          29.83          39          10
## Mean.Wind.SpeedMPH MeanDew.PointF Min.DewpointF Min.Humidity
## 1          13          40          26          52
## 2          13          39          37          89
## 3          13          31          27          82
## 4          11          27          25          64
## 5          12          26          24          55
## 6          10          27          25          53
## Min.Sea.Level.PressureIn Min.TemperatureF Min.VisibilityMiles
## 1          30.01          39          10
## 2          29.43          38           1
## 3          29.44          32           1
## 4          29.81          31           7
## 5          29.78          32          10
## 6          29.78          33          10
## PrecipitationIn WindDirDegrees
## 1          0.01          268
## 2          0.28          357
## 3          0.02          230
## 4           T          286
## 5           T          298
## 6          0.00          306
```

6. Let's look closer at the column types as it is important that variables are coded appropriately for further statistical analysis. This is not yet the case with our weather data. Recall that functions such as `as.numeric()` and `as.character()` can be used to coerce variables into different types.

- Use `str()` to see how variables are stored in `weather5`.
- View the first 20 rows of `weather5`. Keep an eye out for strange values!
- Try coercing the `PrecipitationIn` column of `weather5` to numeric without saving the result.

```
str(weather5) # View the structure of weather5
```

```
## 'data.frame': 366 obs. of 23 variables:
## $ date : Date, format: "2014-12-01" "2014-12-10" ...
## $ Events : chr "Rain" "Rain" "Rain-Snow" "Snow" ...
## $ CloudCover : chr "6" "8" "8" "7" ...
## $ Max.Dew.PointF : chr "46" "45" "37" "28" ...
## $ Max.Gust.SpeedMPH : chr "29" "29" "28" "21" ...
## $ Max.Humidity : chr "74" "100" "92" "85" ...
## $ Max.Sea.Level.PressureIn : chr "30.45" "29.58" "29.81" "29.88" ...
## $ Max.TemperatureF : chr "64" "48" "39" "39" ...
## $ Max.VisibilityMiles : chr "10" "10" "10" "10" ...
## $ Max.Wind.SpeedMPH : chr "22" "23" "21" "16" ...
## $ Mean.Humidity : chr "63" "95" "87" "75" ...
## $ Mean.Sea.Level.PressureIn: chr "30.13" "29.5" "29.61" "29.85" ...
## $ Mean.TemperatureF : chr "52" "43" "36" "35" ...
## $ Mean.VisibilityMiles : chr "10" "3" "7" "10" ...
```

```
## $ Mean.Wind.SpeedMPH      : chr  "13" "13" "13" "11" ...
## $ MeanDew.PointF          : chr  "40" "39" "31" "27" ...
## $ Min.DewpointF           : chr  "26" "37" "27" "25" ...
## $ Min.Humidity            : chr  "52" "89" "82" "64" ...
## $ Min.Sea.Level.PressureIn : chr  "30.01" "29.43" "29.44" "29.81" ...
## $ Min.TemperatureF        : chr  "39" "38" "32" "31" ...
## $ Min.VisibilityMiles     : chr  "10" "1" "1" "7" ...
## $ PrecipitationIn         : chr  "0.01" "0.28" "0.02" "T" ...
## $ WindDirDegrees          : chr  "268" "357" "230" "286" ...
```

```
head(weather5, 20) # Examine the first 20 rows of weather5. Are most of the characters numeric?
```

```
##      date      Events CloudCover Max.Dew.PointF Max.Gust.SpeedMPH
## 1  2014-12-01      Rain           6              46              29
## 2  2014-12-10      Rain           8              45              29
## 3  2014-12-11 Rain-Snow           8              37              28
## 4  2014-12-12      Snow           7              28              21
## 5  2014-12-13              5              28              23
## 6  2014-12-14              4              29              20
## 7  2014-12-15              2              33              21
## 8  2014-12-16      Rain           8              42              10
## 9  2014-12-17      Rain           8              46              26
## 10 2014-12-18      Rain           7              34              30
## 11 2014-12-19              4              25              23
## 12 2014-12-02 Rain-Snow           7              40              29
## 13 2014-12-20      Snow           6              30              26
## 14 2014-12-21      Snow           8              30              20
## 15 2014-12-22      Rain           7              39              22
## 16 2014-12-23      Rain           8              45              25
## 17 2014-12-24 Fog-Rain           8              46              15
## 18 2014-12-25      Rain           6              58              40
## 19 2014-12-26              1              31              25
## 20 2014-12-27              3              34              21
##      Max.Humidity Max.Sea.Level.PressureIn Max.TemperatureF
## 1              74              30.45              64
## 2              100             29.58              48
## 3              92              29.81              39
## 4              85              29.88              39
## 5              75              29.86              42
## 6              82              29.91              45
## 7              89              30.15              42
## 8              96              30.17              44
## 9              100             29.91              49
## 10             89              29.87              44
## 11             69              30.15              37
## 12             92              30.71              42
## 13             89              30.31              36
## 14             85              30.37              36
## 15             89              30.4              44
## 16             100             30.31              47
## 17             100             30.13              46
## 18             100             29.96              59
## 19             70              30.16              50
## 20             70              30.22              52
##      Max.VisibilityMiles Max.Wind.SpeedMPH Mean.Humidity
```

## 1	10	22	63	
## 2	10	23	95	
## 3	10	21	87	
## 4	10	16	75	
## 5	10	17	65	
## 6	10	15	68	
## 7	10	15	75	
## 8	10	8	85	
## 9	10	20	85	
## 10	10	23	73	
## 11	10	17	63	
## 12	10	24	72	
## 13	10	21	79	
## 14	10	16	77	
## 15	10	18	79	
## 16	10	20	91	
## 17	2	13	98	
## 18	10	28	75	
## 19	10	18	60	
## 20	10	17	60	
##	Mean.Sea.Level.PressureIn	Mean.TemperatureF	Mean.VisibilityMiles	
## 1	30.13	52	10	
## 2	29.5	43	3	
## 3	29.61	36	7	
## 4	29.85	35	10	
## 5	29.82	37	10	
## 6	29.83	39	10	
## 7	30.05	37	10	
## 8	30.09	40	9	
## 9	29.75	45	6	
## 10	29.78	40	10	
## 11	29.98	33	10	
## 12	30.59	38	8	
## 13	30.26	32	10	
## 14	30.32	33	9	
## 15	30.35	39	10	
## 16	30.23	45	5	
## 17	29.9	44	1	
## 18	29.63	52	8	
## 19	30.11	44	10	
## 20	30.14	45	10	
##	Mean.Wind.SpeedMPH	MeanDew.PointF	Min.DewpointF	Min.Humidity
## 1	13	40	26	52
## 2	13	39	37	89
## 3	13	31	27	82
## 4	11	27	25	64
## 5	12	26	24	55
## 6	10	27	25	53
## 7	6	29	27	60
## 8	4	36	30	73
## 9	11	41	32	70
## 10	14	30	26	57
## 11	11	22	20	56
## 12	15	27	17	51

```

## 13          10          24          20          69
## 14           9          27          25          69
## 15           8          34          25          69
## 16          13          42          37          82
## 17           6          44          41          96
## 18          14          43          29          49
## 19          11          29          28          49
## 20           9          31          29          50
##   Min.Sea.Level.PressureIn Min.TemperatureF Min.VisibilityMiles
## 1          30.01          39          10
## 2          29.43          38           1
## 3          29.44          32           1
## 4          29.81          31           7
## 5          29.78          32          10
## 6          29.78          33          10
## 7          29.91          32          10
## 8          29.92          35           5
## 9          29.69          41           1
## 10         29.71          36          10
## 11         29.86          29          10
## 12          30.4          33           2
## 13         30.17          27           7
## 14         30.28          30           6
## 15          30.3          33           4
## 16         30.16          42           1
## 17         29.55          41           0
## 18         29.47          44           1
## 19         29.99          37          10
## 20         30.03          38          10
##   PrecipitationIn WindDirDegrees
## 1          0.01          268
## 2          0.28          357
## 3          0.02          230
## 4           T          286
## 5           T          298
## 6          0.00          306
## 7          0.00          324
## 8           T           79
## 9          0.43          311
## 10         0.01          281
## 11         0.00          305
## 12         0.10           62
## 13           T          350
## 14           T           2
## 15         0.05          24
## 16         0.25          63
## 17         0.56          12
## 18         0.14          250
## 19         0.00          255
## 20         0.00          251

```

```
as.numeric(weather5$PrecipitationIn) # See what happens if we try to convert PrecipitationIn to numeric
```

```
## Warning: NAs durch Umwandlung erzeugt
```

```
## [1] 0.01 0.28 0.02 NA NA 0.00 0.00 NA 0.43 0.01 0.00 0.10 NA NA
## [15] 0.05 0.25 0.56 0.14 0.00 0.00 0.01 0.00 0.44 0.00 0.00 0.00 0.11 1.09
## [29] 0.13 0.03 2.90 0.00 0.00 0.00 0.20 0.00 NA 0.12 0.00 0.00 0.15 0.00
## [43] 0.00 0.00 0.00 NA 0.00 0.71 0.00 0.10 0.95 0.01 NA 0.62 0.06 0.05
## [57] 0.57 0.00 0.02 NA 0.00 0.01 0.00 0.05 0.01 0.03 0.00 0.23 0.39 0.00
## [71] 0.02 0.01 0.06 0.78 0.00 0.17 0.11 0.00 NA 0.07 0.02 0.00 0.00 0.00
## [85] 0.00 0.09 NA 0.07 0.37 0.88 0.17 0.06 0.01 0.00 0.00 0.80 0.27 0.00
## [99] 0.14 0.00 0.00 0.01 0.05 0.09 0.00 0.00 0.00 0.04 0.80 0.21 0.12 0.00
## [113] 0.26 NA 0.00 0.02 NA 0.00 0.00 NA 0.00 0.00 0.09 0.00 0.00 0.00
## [127] 0.01 0.00 0.00 0.06 0.00 0.00 0.00 0.61 0.54 NA 0.00 NA 0.00 0.00
## [141] 0.10 0.07 0.00 0.03 0.00 0.39 0.00 0.00 0.03 0.26 0.09 0.00 0.00 0.00
## [155] 0.02 0.00 0.00 0.00 NA 0.00 0.00 0.27 0.00 0.00 0.00 NA 0.00 0.00
## [169] NA 0.00 0.00 NA 0.00 0.00 0.00 0.91 0.00 0.02 0.00 0.00 0.00 0.00
## [183] 0.38 0.00 0.00 0.00 NA 0.00 0.40 NA 0.00 0.00 0.00 0.74 0.04 1.72
## [197] 0.00 0.01 0.00 0.00 NA 0.20 1.43 NA 0.00 0.00 0.00 NA 0.09 0.00
## [211] NA NA 0.50 1.12 0.00 0.00 0.00 0.03 NA 0.00 NA 0.14 NA 0.00
## [225] NA NA 0.00 0.00 0.01 0.00 NA 0.06 0.00 0.00 0.00 0.02 0.00 NA
## [239] 0.00 0.00 0.02 NA 0.15 NA 0.00 0.83 0.00 0.00 0.00 0.08 0.00 0.00
## [253] 0.14 0.00 0.00 0.00 0.63 NA 0.02 NA 0.00 NA 0.00 0.00 0.00 0.00
## [267] 0.00 0.00 0.49 0.00 0.00 0.00 0.00 0.00 0.00 0.17 0.66 0.01 0.38 0.00
## [281] 0.00 0.00 0.00 0.00 0.00 0.00 NA 0.00 0.00 0.00 0.00 0.00 0.00 0.00
## [295] 0.00 0.04 0.01 2.46 NA 0.00 0.00 0.00 0.20 0.00 NA 0.00 0.00 0.00
## [309] 0.12 0.00 0.00 NA NA NA 0.00 0.08 NA 0.07 NA 0.00 0.00 0.03
## [323] 0.00 0.00 0.36 0.73 0.01 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.34 NA
## [337] 0.07 0.54 0.04 0.01 0.00 0.00 0.00 0.00 0.00 0.00 NA 0.00 0.86 0.00 0.30
## [351] 0.04 0.00 0.00 0.00 0.00 0.21 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00
## [365] 0.00 0.14
```

Scroll the output, notice the warning message. Go back to the results of the head command if need be. What values in PrecipitationIn would become NA if coerced to numbers? Why would they be in the dataset to begin with?

7. As you saw in the last exercise, T was used to denote a trace amount (i.e. too small to be accurately measured) of precipitation in the PrecipitationIn column. In order to coerce this column to numeric, you'll need to deal with this somehow. To keep things simple, we will just replace T with zero, as a string ("0").
- Use `str_replace()` from `stringr` to make the proper replacements in the PrecipitationIn column of `weather5`.
- Run the call to `mutate_at` to conveniently apply `as.numeric()` to all columns from CloudCover through WindDirDegrees (reading left to right in the data), saving the result to `weather6`.
- View the structure of `weather6` to confirm the coercions were successful.

```
weather5$PrecipitationIn <- str_replace(weather5$PrecipitationIn, "T", "0") # Replace "T" with "0" (T =
weather6 <- mutate_at(weather5, vars(CloudCover:WindDirDegrees), funs(as.numeric)) # Convert characters
```

```
## Warning: funs() is soft deprecated as of dplyr 0.8.0
## please use list() instead
##
## # Before:
## funs(name = f())
##
## # After:
## list(name = ~f())
## This warning is displayed once per session.
```

```
str(weather6) # Look at result
```

```
## 'data.frame':   366 obs. of  23 variables:
## $ date          : Date, format: "2014-12-01" "2014-12-10" ...
## $ Events        : chr  "Rain" "Rain" "Rain-Snow" "Snow" ...
## $ CloudCover     : num  6 8 8 7 5 4 2 8 8 7 ...
## $ Max.Dew.PointF : num  46 45 37 28 28 29 33 42 46 34 ...
## $ Max.Gust.SpeedMPH : num  29 29 28 21 23 20 21 10 26 30 ...
## $ Max.Humidity    : num  74 100 92 85 75 82 89 96 100 89 ...
## $ Max.Sea.Level.PressureIn : num  30.4 29.6 29.8 29.9 29.9 ...
## $ Max.TemperatureF : num  64 48 39 39 42 45 42 44 49 44 ...
## $ Max.VisibilityMiles : num  10 10 10 10 10 10 10 10 10 10 ...
## $ Max.Wind.SpeedMPH : num  22 23 21 16 17 15 15 8 20 23 ...
## $ Mean.Humidity    : num  63 95 87 75 65 68 75 85 85 73 ...
## $ Mean.Sea.Level.PressureIn: num  30.1 29.5 29.6 29.9 29.8 ...
## $ Mean.TemperatureF : num  52 43 36 35 37 39 37 40 45 40 ...
## $ Mean.VisibilityMiles : num  10 3 7 10 10 10 10 9 6 10 ...
## $ Mean.Wind.SpeedMPH : num  13 13 13 11 12 10 6 4 11 14 ...
## $ MeanDew.PointF   : num  40 39 31 27 26 27 29 36 41 30 ...
## $ Min.DewpointF    : num  26 37 27 25 24 25 27 30 32 26 ...
## $ Min.Humidity      : num  52 89 82 64 55 53 60 73 70 57 ...
## $ Min.Sea.Level.PressureIn : num  30 29.4 29.4 29.8 29.8 ...
## $ Min.TemperatureF : num  39 38 32 31 32 33 32 35 41 36 ...
## $ Min.VisibilityMiles : num  10 1 1 7 10 10 10 5 1 10 ...
## $ PrecipitationIn  : num  0.01 0.28 0.02 0 0 0 0 0 0.43 0.01 ...
## $ WindDirDegrees   : num  268 357 230 286 298 306 324 79 311 281 ...
```

8. Before dealing with missing values in the data, it's important to find them and figure out why they exist in the first place. If your dataset is too big to look at all at once, like it is here, remember you can use `sum()` and `is.na()` to quickly size up the situation by counting the number of NA values. The `summary()` function may also come in handy for identifying which variables contain the missing values. Finally, the `which()` function is useful for locating the missing values within a particular column.

- Use `sum()` and `is.na()` to count the number of NA values in `weather6`.
- Look at a `summary()` of `weather6` to figure out how the missings are distributed among the different variables.
- Use `which()` to identify the indices (i.e. row numbers) where `Max.Gust.SpeedMPH` is NA and save the result to `ind`.
- Use `ind` to look at the full rows of `weather6` for which `Max.Gust.SpeedMPH` is missing.

```
sum(is.na(weather6)) # Count missing values
```

```
## [1] 6
```

```
summary(weather6) # Find missing values
```

```
##      date          Events          CloudCover    Max.Dew.PointF
## Min.   :2014-12-01  Length:366      Min.     :0.000    Min.     : -6.00
## 1st Qu.:2015-03-02  Class :character  1st Qu.:3.000    1st Qu.:32.00
## Median :2015-06-01  Mode  :character  Median :5.000    Median :47.50
## Mean   :2015-06-01                Mean   :4.708    Mean   :45.48
## 3rd Qu.:2015-08-31                3rd Qu.:7.000    3rd Qu.:61.00
## Max.   :2015-12-01                Max.    :8.000    Max.    :75.00
##
## Max.Gust.SpeedMPH  Max.Humidity    Max.Sea.Level.PressureIn
```



```

## Min. : 0.00      Min. : 39.00      Min. :29.58
## 1st Qu.:21.00     1st Qu.: 73.25     1st Qu.:30.00
## Median :25.50     Median : 86.00     Median :30.14
## Mean :26.99       Mean : 85.69       Mean :30.16
## 3rd Qu.:31.25     3rd Qu.: 93.00     3rd Qu.:30.31
## Max. :94.00       Max. :1000.00      Max. :30.88
## NA's :6
## Max.TemperatureF Max.VisibilityMiles Max.Wind.SpeedMPH Mean.Humidity
## Min. :18.00      Min. : 2.000      Min. : 8.00      Min. :28.00
## 1st Qu.:42.00     1st Qu.:10.000     1st Qu.:16.00     1st Qu.:56.00
## Median :60.00     Median :10.000     Median :20.00     Median :66.00
## Mean :58.93       Mean : 9.907       Mean :20.62       Mean :66.02
## 3rd Qu.:76.00     3rd Qu.:10.000     3rd Qu.:24.00     3rd Qu.:76.75
## Max. :96.00       Max. :10.000       Max. :38.00       Max. :98.00
##
## Mean.Sea.Level.PressureIn Mean.TemperatureF Mean.VisibilityMiles
## Min. :29.49      Min. : 8.00      Min. : -1.000
## 1st Qu.:29.87     1st Qu.:36.25     1st Qu.: 8.000
## Median :30.03     Median :53.50     Median :10.000
## Mean :30.04       Mean :51.40       Mean : 8.861
## 3rd Qu.:30.19     3rd Qu.:68.00     3rd Qu.:10.000
## Max. :30.77       Max. :84.00       Max. :10.000
##
## Mean.Wind.SpeedMPH MeanDew.PointF Min.DewpointF Min.Humidity
## Min. : 4.00      Min. : -11.00     Min. : -18.00     Min. :16.00
## 1st Qu.: 8.00     1st Qu.: 24.00     1st Qu.: 16.25     1st Qu.:35.00
## Median :10.00     Median : 41.00     Median : 35.00     Median :46.00
## Mean :10.68       Mean : 38.96       Mean : 32.25       Mean :48.31
## 3rd Qu.:13.00     3rd Qu.: 56.00     3rd Qu.: 51.00     3rd Qu.:60.00
## Max. :22.00       Max. : 71.00       Max. : 68.00       Max. :96.00
##
## Min.Sea.Level.PressureIn Min.TemperatureF Min.VisibilityMiles
## Min. :29.16      Min. : -3.00      Min. : 0.000
## 1st Qu.:29.76     1st Qu.:30.00     1st Qu.: 2.000
## Median :29.94     Median :46.00     Median :10.000
## Mean :29.93       Mean :43.33       Mean : 6.716
## 3rd Qu.:30.09     3rd Qu.:60.00     3rd Qu.:10.000
## Max. :30.64       Max. :74.00       Max. :10.000
##
## PrecipitationIn WindDirDegrees
## Min. :0.0000      Min. : 1.0
## 1st Qu.:0.0000     1st Qu.:113.0
## Median :0.0000     Median :222.0
## Mean :0.1016       Mean :200.1
## 3rd Qu.:0.0400     3rd Qu.:275.0
## Max. :2.9000       Max. :360.0
##

```

```

ind <- which(is.na(weather6$Max.Gust.SpeedMPH)) # Find indices of NAs in Max.Gust.SpeedMPH
weather6[ind, ] # Look at the full rows for records missing Max.Gust.SpeedMPH

```

```

##           date Events CloudCover Max.Dew.PointF Max.Gust.SpeedMPH
## 161 2015-05-18   Fog           6           52           NA
## 205 2015-06-03           7           48           NA
## 273 2015-08-08           4           61           NA

```

```
## 275 2015-09-01          1          63          NA
## 308 2015-10-12          0          56          NA
## 358 2015-11-03          1          44          NA
##      Max.Humidity Max.Sea.Level.PressureIn Max.TemperatureF
## 161          100          30.30          58
## 205           93          30.31          56
## 273           87          30.02          76
## 275           78          30.06          79
## 308           89          29.86          76
## 358           82          30.25          73
##      Max.VisibilityMiles Max.Wind.SpeedMPH Mean.Humidity
## 161           10          16          79
## 205           10          14          82
## 273           10          14          68
## 275           10          15          65
## 308           10          15          65
## 358           10          16          57
##      Mean.Sea.Level.PressureIn Mean.TemperatureF Mean.VisibilityMiles
## 161          30.23          54           8
## 205          30.24          52          10
## 273          29.99          69          10
## 275          30.02          74          10
## 308          29.80          64          10
## 358          30.13          60          10
##      Mean.Wind.SpeedMPH MeanDew.PointF Min.DewpointF Min.Humidity
## 161           10          48          43          57
## 205            7          45          43          71
## 273            6          57          54          49
## 275            9          62          59          52
## 308            8          51          48          41
## 358            8          42          40          31
##      Min.Sea.Level.PressureIn Min.TemperatureF Min.VisibilityMiles
## 161          30.12          49           0
## 205          30.19          47          10
## 273          29.95          61          10
## 275          29.96          69          10
## 308          29.74          51          10
## 358          30.06          47          10
##      PrecipitationIn WindDirDegrees
## 161            0          72
## 205            0          90
## 273            0          45
## 275            0          54
## 308            0         199
## 358            0         281
```

9. Besides missing values, we want to know if there are values in the data that are too extreme or bizarre to be plausible. A great way to start the search for these values is with `summary()`. Once implausible values are identified, they must be dealt with in an intelligent and informed way. Sometimes the best way forward is obvious and other times it may require some research and/or discussions with the original collectors of the data.

- View a `summary()` of `weather6`.
- Use `which()` to find the index of the erroneous element of `weather6$Max.Humidity`, saving the result

to ind.

- Use ind to look at the full row of `weather6` for that day. You discover an extra zero was accidentally added to this value. Correct it in the data.

```
summary(weather6) # Review distributions for all variables
```

```
##      date      Events      CloudCover      Max.Dew.PointF
## Min.   :2014-12-01  Length:366      Min.    :0.000      Min.    :-6.00
## 1st Qu.:2015-03-02  Class :character  1st Qu.:3.000      1st Qu.:32.00
## Median :2015-06-01  Mode  :character  Median :5.000      Median :47.50
## Mean   :2015-06-01                      Mean   :4.708      Mean   :45.48
## 3rd Qu.:2015-08-31                      3rd Qu.:7.000      3rd Qu.:61.00
## Max.   :2015-12-01                      Max.    :8.000      Max.    :75.00
##
## Max.Gust.SpeedMPH  Max.Humidity      Max.Sea.Level.PressureIn
## Min.   : 0.00      Min.    : 39.00      Min.    :29.58
## 1st Qu.:21.00      1st Qu.: 73.25      1st Qu.:30.00
## Median :25.50      Median : 86.00      Median :30.14
## Mean   :26.99      Mean   : 85.69      Mean   :30.16
## 3rd Qu.:31.25      3rd Qu.: 93.00      3rd Qu.:30.31
## Max.   :94.00      Max.    :1000.00     Max.    :30.88
## NA's      :6
## Max.TemperatureF  Max.VisibilityMiles  Max.Wind.SpeedMPH  Mean.Humidity
## Min.   :18.00      Min.    : 2.000      Min.    : 8.00      Min.    :28.00
## 1st Qu.:42.00      1st Qu.:10.000      1st Qu.:16.00      1st Qu.:56.00
## Median :60.00      Median :10.000      Median :20.00      Median :66.00
## Mean   :58.93      Mean   : 9.907      Mean   :20.62      Mean   :66.02
## 3rd Qu.:76.00      3rd Qu.:10.000      3rd Qu.:24.00      3rd Qu.:76.75
## Max.   :96.00      Max.    :10.000      Max.    :38.00      Max.    :98.00
##
## Mean.Sea.Level.PressureIn  Mean.TemperatureF  Mean.VisibilityMiles
## Min.   :29.49              Min.    : 8.00      Min.    : -1.000
## 1st Qu.:29.87              1st Qu.:36.25      1st Qu.: 8.000
## Median :30.03              Median :53.50      Median :10.000
## Mean   :30.04              Mean   :51.40      Mean   : 8.861
## 3rd Qu.:30.19              3rd Qu.:68.00      3rd Qu.:10.000
## Max.   :30.77              Max.    :84.00      Max.    :10.000
##
## Mean.Wind.SpeedMPH  MeanDew.PointF      Min.DewpointF      Min.Humidity
## Min.   : 4.00      Min.    : -11.00     Min.    : -18.00     Min.    :16.00
## 1st Qu.: 8.00      1st Qu.: 24.00     1st Qu.: 16.25     1st Qu.:35.00
## Median :10.00      Median : 41.00     Median : 35.00     Median :46.00
## Mean   :10.68      Mean   : 38.96     Mean   : 32.25     Mean   :48.31
## 3rd Qu.:13.00      3rd Qu.: 56.00     3rd Qu.: 51.00     3rd Qu.:60.00
## Max.   :22.00      Max.    : 71.00     Max.    : 68.00     Max.    :96.00
##
## Min.Sea.Level.PressureIn  Min.TemperatureF  Min.VisibilityMiles
## Min.   :29.16              Min.    : -3.00     Min.    : 0.000
## 1st Qu.:29.76              1st Qu.:30.00     1st Qu.: 2.000
## Median :29.94              Median :46.00     Median :10.000
## Mean   :29.93              Mean   :43.33     Mean   : 6.716
## 3rd Qu.:30.09              3rd Qu.:60.00     3rd Qu.:10.000
## Max.   :30.64              Max.    :74.00     Max.    :10.000
##
```

```
## PrecipitationIn WindDirDegrees
## Min. :0.0000 Min. : 1.0
## 1st Qu.:0.0000 1st Qu.:113.0
## Median :0.0000 Median :222.0
## Mean :0.1016 Mean :200.1
## 3rd Qu.:0.0400 3rd Qu.:275.0
## Max. :2.9000 Max. :360.0
##
```

```
ind <- which(weather6$Max.Humidity == 1000) # Find row with Max.Humidity of 1000
weather6[ind, ] # Look at the data for that day
```

```
##          date          Events CloudCover Max.Dew.PointF
## 135 2015-04-21 Fog-Rain-Thunderstorm          6          57
##      Max.Gust.SpeedMPH Max.Humidity Max.Sea.Level.PressureIn
## 135          94          1000          29.75
##      Max.TemperatureF Max.VisibilityMiles Max.Wind.SpeedMPH Mean.Humidity
## 135          65          10          20          71
##      Mean.Sea.Level.PressureIn Mean.TemperatureF Mean.VisibilityMiles
## 135          29.6          56          5
##      Mean.Wind.SpeedMPH MeanDew.PointF Min.DewpointF Min.Humidity
## 135          10          49          36          42
##      Min.Sea.Level.PressureIn Min.TemperatureF Min.VisibilityMiles
## 135          29.53          46          0
##      PrecipitationIn WindDirDegrees
## 135          0.54          184
```

```
weather6$Max.Humidity[ind] <- 100 # Change 1000 to 100
```

10. You've discovered and repaired one obvious error in the data, but it appears that there's another. Sometimes you get lucky and can infer the correct or intended value from the other data. For example, if you know the minimum and maximum values of a particular metric on a given day.

- Use `summary()` to look at the value of only the `Mean.VisibilityMiles` variable of `weather6`.
- Determine the element of the value that is clearly erroneous in this column, saving the result to `ind`.
- Use `ind` to look at the full row of `weather6` for this day.
- Inspect the values of other variables for this day to determine the correct value of `Mean.VisibilityMiles`, then make the appropriate fix.

```
summary(weather6$Mean.VisibilityMiles) # Look at summary of Mean.VisibilityMiles
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -1.000  8.000 10.000  8.861 10.000 10.000
```

```
ind <- which(weather6$Mean.VisibilityMiles == -1) # Get index of row with -1 value
weather6[ind, ] # Look at full row
```

```
##          date Events CloudCover Max.Dew.PointF Max.Gust.SpeedMPH
## 192 2015-06-18          5          54          23
##      Max.Humidity Max.Sea.Level.PressureIn Max.TemperatureF
## 192          72          30.14          76
##      Max.VisibilityMiles Max.Wind.SpeedMPH Mean.Humidity
## 192          10          17          59
##      Mean.Sea.Level.PressureIn Mean.TemperatureF Mean.VisibilityMiles
## 192          30.04          67          -1
##      Mean.Wind.SpeedMPH MeanDew.PointF Min.DewpointF Min.Humidity
```

```
## 192          10          49          45          46
##   Min.Sea.Level.PressureIn Min.TemperatureF Min.VisibilityMiles
## 192          29.93          57          10
##   PrecipitationIn WindDirDegrees
## 192          0          189
weather6$Mean.VisibilityMiles[ind] <- 10 # Set Mean.VisibilityMiles to the appropriate value
```

11. In addition to dealing with obvious errors in the data, we want to see if there are other extreme values. In addition to the trusty `summary()` function, `hist()` is useful for quickly getting a feel for how different variables are distributed.

- Check a `summary()` of `weather6` one more time for extreme or unexpected values.
- View a histogram for `MeanDew.PointF`, `Min.TemperatureF` and `Mean.TemperatureF` to compare distributions.

```
summary(weather6) # Review summary of full data once more
```

```
##      date          Events      CloudCover      Max.Dew.PointF
## Min.   :2014-12-01  Length:366      Min.    :0.000      Min.    :-6.00
## 1st Qu.:2015-03-02  Class :character  1st Qu.:3.000      1st Qu.:32.00
## Median :2015-06-01  Mode  :character  Median :5.000      Median :47.50
## Mean   :2015-06-01                      Mean  :4.708      Mean  :45.48
## 3rd Qu.:2015-08-31                      3rd Qu.:7.000      3rd Qu.:61.00
## Max.   :2015-12-01                      Max.   :8.000      Max.   :75.00
##
## Max.Gust.SpeedMPH  Max.Humidity  Max.Sea.Level.PressureIn
## Min.    : 0.00      Min.    : 39.00      Min.    :29.58
## 1st Qu.:21.00      1st Qu.: 73.25      1st Qu.:30.00
## Median :25.50      Median : 86.00      Median :30.14
## Mean   :26.99      Mean   : 83.23      Mean   :30.16
## 3rd Qu.:31.25      3rd Qu.: 93.00      3rd Qu.:30.31
## Max.   :94.00      Max.   :100.00      Max.   :30.88
## NA's    :6
## Max.TemperatureF  Max.VisibilityMiles  Max.Wind.SpeedMPH  Mean.Humidity
## Min.    :18.00      Min.    : 2.000      Min.    : 8.00      Min.    :28.00
## 1st Qu.:42.00      1st Qu.:10.000      1st Qu.:16.00      1st Qu.:56.00
## Median :60.00      Median :10.000      Median :20.00      Median :66.00
## Mean   :58.93      Mean   : 9.907      Mean   :20.62      Mean   :66.02
## 3rd Qu.:76.00      3rd Qu.:10.000      3rd Qu.:24.00      3rd Qu.:76.75
## Max.   :96.00      Max.   :10.000      Max.   :38.00      Max.   :98.00
##
## Mean.Sea.Level.PressureIn  Mean.TemperatureF  Mean.VisibilityMiles
## Min.    :29.49              Min.    : 8.00      Min.    : 1.000
## 1st Qu.:29.87              1st Qu.:36.25      1st Qu.: 8.000
## Median :30.03              Median :53.50      Median :10.000
## Mean   :30.04              Mean   :51.40      Mean   : 8.891
## 3rd Qu.:30.19              3rd Qu.:68.00      3rd Qu.:10.000
## Max.   :30.77              Max.   :84.00      Max.   :10.000
##
## Mean.Wind.SpeedMPH  MeanDew.PointF  Min.DewpointF  Min.Humidity
## Min.    : 4.00      Min.    : -11.00  Min.    : -18.00  Min.    :16.00
## 1st Qu.: 8.00      1st Qu.: 24.00  1st Qu.: 16.25  1st Qu.:35.00
## Median :10.00      Median : 41.00  Median : 35.00  Median :46.00
## Mean   :10.68      Mean   : 38.96  Mean   : 32.25  Mean   :48.31
```

```
## 3rd Qu.:13.00      3rd Qu.: 56.00   3rd Qu.: 51.00   3rd Qu.:60.00
## Max.    :22.00      Max.     : 71.00   Max.     : 68.00   Max.     :96.00
##
## Min.Sea.Level.PressureIn Min.TemperatureF Min.VisibilityMiles
## Min.    :29.16      Min.     :-3.00   Min.     : 0.000
## 1st Qu.:29.76      1st Qu.:30.00   1st Qu.: 2.000
## Median :29.94      Median :46.00   Median :10.000
## Mean    :29.93      Mean    :43.33   Mean     : 6.716
## 3rd Qu.:30.09      3rd Qu.:60.00   3rd Qu.:10.000
## Max.    :30.64      Max.     :74.00   Max.     :10.000
##
## PrecipitationIn WindDirDegrees
## Min.     :0.0000   Min.      : 1.0
## 1st Qu.:0.0000   1st Qu.:113.0
## Median :0.0000   Median :222.0
## Mean    :0.1016   Mean     :200.1
## 3rd Qu.:0.0400   3rd Qu.:275.0
## Max.    :2.9000   Max.     :360.0
##
```

```
hist(weather6$MeanDew.PointF) # Look at histogram for MeanDew.PointF
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-142-1.pdf

```
hist(weather6$Min.TemperatureF) # Look at histogram for Min.TemperatureF
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-142-2.pdf

```
hist(weather6$Mean.TemperatureF) # Compare to histogram for Mean.TemperatureF
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-142-3.pdf

12. Finally, the `Events` column contains an empty string (“”) for any day on which there was no significant weather event such as rain, fog, a thunderstorm, etc. However, if it’s the first time you’re seeing these data, it may not be obvious that this is the case, so it’s best for us to be explicit and replace the empty strings with something more meaningful. Replace all empty strings in the `events` column of `weather6` with “None”. One last time, print out the first 6 rows of the `weather6` data frame to see the changes.

```
weather6$Events[weather6$Events == ""] <- "None" # Replace empty cells in events column
head(weather6) # Print the first 6 rows of weather6
```

```
##           date      Events CloudCover Max.Dew.PointF Max.Gust.SpeedMPH
```

## 1	2014-12-01	Rain	6	46	29
## 2	2014-12-10	Rain	8	45	29
## 3	2014-12-11	Rain-Snow	8	37	28
## 4	2014-12-12	Snow	7	28	21
## 5	2014-12-13	None	5	28	23
## 6	2014-12-14	None	4	29	20
##	Max.Humidity	Max.Sea.Level.PressureIn	Max.TemperatureF		
## 1	74	30.45	64		
## 2	100	29.58	48		
## 3	92	29.81	39		
## 4	85	29.88	39		
## 5	75	29.86	42		
## 6	82	29.91	45		
##	Max.VisibilityMiles	Max.Wind.SpeedMPH	Mean.Humidity		
## 1	10	22	63		
## 2	10	23	95		
## 3	10	21	87		
## 4	10	16	75		
## 5	10	17	65		
## 6	10	15	68		
##	Mean.Sea.Level.PressureIn	Mean.TemperatureF	Mean.VisibilityMiles		
## 1	30.13	52	10		
## 2	29.50	43	3		
## 3	29.61	36	7		
## 4	29.85	35	10		
## 5	29.82	37	10		
## 6	29.83	39	10		
##	Mean.Wind.SpeedMPH	MeanDew.PointF	Min.DewpointF	Min.Humidity	
## 1	13	40	26	52	
## 2	13	39	37	89	
## 3	13	31	27	82	
## 4	11	27	25	64	
## 5	12	26	24	55	
## 6	10	27	25	53	
##	Min.Sea.Level.PressureIn	Min.TemperatureF	Min.VisibilityMiles		
## 1	30.01	39	10		
## 2	29.43	38	1		
## 3	29.44	32	1		
## 4	29.81	31	7		
## 5	29.78	32	10		
## 6	29.78	33	10		
##	PrecipitationIn	WindDirDegrees			
## 1	0.01	268			
## 2	0.28	357			
## 3	0.02	230			
## 4	0.00	286			
## 5	0.00	298			
## 6	0.00	306			

User-defined functions

1. Define a Cobb-Douglas production function with two inputs vectors,

$$x = \begin{pmatrix} L \\ K \end{pmatrix}$$
$$\theta = \begin{pmatrix} A \\ \alpha \\ \beta \end{pmatrix}$$

and scalar output

$$y = AL^\alpha K^\beta.$$

Evaluate the function at

$$x = \begin{pmatrix} 2 \\ 3 \end{pmatrix}$$
$$\theta = \begin{pmatrix} 1 \\ 0.3 \\ 0.8 \end{pmatrix}.$$

```
cobb_douglas <- function(x, theta) {  
  y <- theta[1] * x[1]^theta[2] * x[2]^theta[3]  
  return(y)  
}  
cobb_douglas(x = c(2, 3), theta = c(1, 0.3, 0.8))
```

```
## [1] 2.964872
```

2. Define a function `lowdecile` with one input vector (x_1, \dots, x_n) of arbitrary length. The function should compute and return the mean of all observations in the lowest decile. Define the vector

$$x = (0, 0, 0, 0, 1, 1, 1, 1, 2, 2, 2, \dots, 9, 9, 9, 9)$$

and apply `lowdecile` to x .

```
lowdecile <- function(x) {  
  quantil <- x[x <= quantile(x, p = 0.1)]  
  return(mean(quantil))  
}  
lowdecile(x = rep(0:9, each = 4))
```

```
## [1] 0
```

Programming

1. This exercise illustrates that loops are often not very efficient.
 - Create the vector $x = (1, 2, \dots, 1\,000\,000)$ and convert it from *integer* to *numeric* using the conversion command `as.numeric`.


```
x <- 1:1e+06
class(x)
```

```
## [1] "integer"
```

```
x <- as.numeric(x)
class(x)
```

```
## [1] "numeric"
```

- Write a for-loop to compute the sum of all vector elements without using the `sum` command. Put the command `p0 <- proc.time()[3]` in front of the loop and the command `print(proc.time()[3]-p0)` at the end. These commands allow to measure the execution time of the loop.

```
S <- 0 # initialize
p0 <- proc.time()[3] #Startzeitpunkt festlegen
for (i in x) {
  S <- S + i
}
print(S)
```

```
## [1] 500000500000
```

```
print(proc.time()[3] - p0) #time used
```

```
## elapsed
## 0.071
```

*Compare your result with the execution time of the `sum` command.

```
p0 <- proc.time()[3]
sum(x)
```

```
## [1] 500000500000
```

```
print(proc.time()[3] - p0)
```

```
## elapsed
## 0.001
```

2. Create a grid vector x of 60 equidistant points x_1, \dots, x_{60} on the interval $[-10, 10]$, and another grid vector y of 70 points y_1, \dots, y_{70} on $[-10, 10]$. Create an empty matrix Z of dimension 60×70 .

Write a double loop to compute the matrix elements

$$Z_{ij} = \frac{10}{r_{ij}} \cdot \sin(r_{ij})$$

where $r_{ij} = \sqrt{x_i^2 + y_j^2}$. Execute `persp(x,y,Z)`.

```
x <- seq(-10, 10, length = 60)
y <- seq(-10, 10, length = 70)
Z <- matrix(NA, 60, 70)
for (i in 1:length(x)) {
  for (j in 1:length(y)) {
    r <- sqrt(x[i]^2 + y[j]^2) # note that r is overwritten in each run of the loop
    Z[i, j] <- 10/r * sin(r)
  }
}
```

```
}
persp(x, y, Z, ticktype = "detailed", col = "lightblue")
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-151-1.pdf

3. Load the data set **fussballdaten.csv**. It contains all 1. *Bundesliga* results between the seasons 1996/1997 and 2008/2009.

```
fussballdaten <- read.csv2("data/fussballdaten.csv", as.is = TRUE)
```

- Create an alphabetically ordered vector of all clubs in the data set.

```
home <- fussballdaten$Heim
away <- fussballdaten$Auswaerts
clubs <- sort(unique(home))
```

- Write a loop over all clubs. For each club compute the proportion of games won.

```
ngames <- dim(fussballdaten)[1] # number of games in dataset
GoalsH <- fussballdaten$ToreH
GoalsA <- fussballdaten$ToreA
winner <- rep(NA, ngames)
propwin <- rep(NA, length(clubs))
# Get winning teams
for (i in 1:ngames) {
  if (GoalsH[i] > GoalsA[i]) {
    winner[i] <- home[i]
  }
  if (GoalsA[i] > GoalsH[i]) {
    winner[i] <- away[i]
  }
  if (GoalsH[i] == GoalsA[i]) {
    winner[i] <- "Remis"
  }
}
# Get proportions
for (i in 1:length(clubs)) {
  win <- sum(winner == clubs[i]) # Games won by club i
  tot <- sum(home == clubs[i] | away == clubs[i]) # Number of matches of club i
  propwin[i] <- win/tot
}
names(propwin) <- clubs
print(propwin)
```

```
## 1860muenchen      aachen      bielefeld      bmuennen      bochum
##      0.3496732    0.2647059    0.2720588      0.6153846    0.3202614
##      bremen      cottbus      dortmund      duesseldorf    duisburg
##      0.4751131    0.2823529    0.4298643      0.2500000    0.2598039
##      fckoeln     frankfurt     freiburg      hamburg      hannover
##      0.2689076    0.2830882    0.2720588      0.3891403    0.3235294
##      herthabsc   karlsruhe    klautern      leverkusen     mainz
```


```
##      0.4144385      0.3308824      0.3823529      0.4728507      0.2843137
##      mgladbach      nuernberg      rostock      schalke      stpauli
##      0.2764706      0.2731092      0.3048128      0.4343891      0.1960784
##      uerdingen      uhaching      ulm      vfbstuttgart      wolfsburg
##      0.1470588      0.2941176      0.2647059      0.4185520      0.3582888
```

- Order the clubs descendingly according to the proportion of games won and plot a **barplot** of the proportion.

```
sort(propwin, decreasing = TRUE)
```

```
##      bmuennen      bremen      leverkusen      schalke      dortmund
##      0.6153846      0.4751131      0.4728507      0.4343891      0.4298643
##      vfbstuttgart      herthabsc      hamburg      klautern      wolfsburg
##      0.4185520      0.4144385      0.3891403      0.3823529      0.3582888
##      1860muenchen      karlsruhe      hannover      bochum      rostock
##      0.3496732      0.3308824      0.3235294      0.3202614      0.3048128
##      uhaching      mainz      frankfurt      cottbus      mgladbach
##      0.2941176      0.2843137      0.2830882      0.2823529      0.2764706
##      nuernberg      bielefeld      freiburg      fckoeln      aachen
##      0.2731092      0.2720588      0.2720588      0.2689076      0.2647059
##      ulm      duisburg      duesseldorf      stpauli      uerdingen
##      0.2647059      0.2598039      0.2500000      0.1960784      0.1470588
```

```
barplot(sort(propwin, decreasing = TRUE), col = "steelblue", las = 3) # label is printed vertically us
```



Random numbers

This section is not only about random number generation but also includes exercises about the R-functions for standard distributions in statistics.

1. Let's consider a simple count data example.

- Let $X \sim N(0, 1)$. Compute the probability $P(|X| > 3.5)$.

```
2 * (1 - pnorm(3.5))
```

```
## [1] 0.0004652582
```

- Generate $n = 10000$ random draws X_1, \dots, X_n from X and count the number of observations $|X_i| > 3.5$.

```
n <- 10000
X <- rnorm(n)
sum(abs(X) > 3.5)
```

```
## [1] 4
```

```
sum(abs(X) > 3.5)/n
```

```
## [1] 4e-04
```

the larger n the closer it is to the theoretical value of 0.0004652582

- Repeat drawing random samples $R = 5000$ times and write the counts into a vector Z_1, \dots, Z_{5000} of length 5000.

```
R <- 5000
Z <- rep(NA, R)
n <- 10000
for (i in 1:R) {
  X <- rnorm(n)
  Z[i] <- sum(abs(X) > 3.5)
}
```

- Tabulate Z and compare the frequencies with the probability function of a suitably fitted Poisson distribution.

```
table(Z)
```

```
## Z
##  0  1  2  3  4  5  6  7  8  9 10 11 12 13
## 57 245 466 787 943 867 661 448 280 125 77 31 10 3
```

```
t(data.frame(observ = 0:16, prob = dpois(0:16, lambda=mean(Z))*R))
```

```
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## observ 0.00000 1.0000 2.0000 3.0000 4.0000 5.0000 6.0000
## prob  46.56239 217.7444 509.1299 793.6317 927.8348 867.7853 676.3519
##      [,8]      [,9]     [,10]     [,11]     [,12]     [,13]     [,14]
## observ 7.0000 8.0000 9.0000 10.0000 11.0000 12.0000 13.0000
## prob  451.8417 264.1241 137.2389 64.17838 27.28398 10.63257 3.82478
##      [,15]      [,16]      [,17]
## observ 14.00000 15.00000 16.00000
## prob   1.277586 0.3983001 0.1164132
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'


## The following object is masked from 'package:dplyr':
##
##      select
```

```
truehist(Z, prob = F)
x <- seq(0, 16)
lines(x, dpois(x, lambda = mean(Z)) * R, lwd = 2)
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-160-1.pdf

2. Generate $n = 10000$ draws from a log normal distribution $X \sim e^Y$ where $Y \sim N(1, 0.5^2)$ (the parameters in the R function are `meanlog=1` and `sdlog=0.5`). Split the screen into two plotting areas using the command `par(mfrow=c(2,1))`. Plot the histograms of X and $\ln X$.

```
n <- 10000
x <- rlnorm(n, meanlog = 1, sdlog = 0.5)
par(mfrow = c(2, 1))
truehist(x)
truehist(log(x))
```




IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-161-1.pdf

3. Generate $n = 10000$ draws from $X \sim N(0, 1)$. Compute the cumulated means, i.e.

$$\bar{X}_j = \frac{1}{j} \sum_{i=1}^j X_i$$

for $j = 1, \dots, n$ and plot them. Hint: Use the command `cumsum`.

```
par(mfrow = c(1, 1))
n <- 10000
X <- rnorm(n)
m <- rep(NA, n)
for (i in 1:n) {
  m[i] <- cumsum(X)[i]/i
}
plot(m)
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-162-1.pdf

Simulations

- This exercise illustrates the one-sample t -test.
 - Generate $n = 10$ observations from $X \sim N(10, 3^2)$. Compute the mean and the standard deviation of X_1, \dots, X_{10} .

```
n <- 10
X <- rnorm(n, mean = 10, sd = 3)
m <- mean(X)
s <- sd(X)
print(m)
```

```
## [1] 10.08052
```

```
print(s)
```

```
## [1] 3.344536
```

- The t -statistics of the hypothesis test $H_0 : \mu = 10$ against $H_1 : \mu \neq 10$ is

$$t = \sqrt{10} \frac{\bar{X} - 10}{sd}$$

where sd is the standard deviation (as computed by `sd`). Compute the t -statistic.

```
t <- sqrt(n) * (m - 10)/s
print(t)
```

```
## [1] 0.07613527
```

- Create an empty vector Z of length $R = 5000$. Write a loop over $r = 1, \dots, R$ and repeat the above steps for each r . Save the t -statistic at Z_r .

```
R <- 5000
Z <- rep(NA, R)
for (r in 1:R) {
  X <- rnorm(n, mean = 10, sd = 3)
  m <- mean(X)
  s <- sd(X)
  Z[r] <- sqrt(n) * (m - 10)/s
}
```

- Plot the histogram of Z_1, \dots, Z_R and add the density function of the t_9 -distribution.

```
library(MASS)
truehist(Z, col = "lightblue")
x <- seq(-4, 4, by = 0.1)
lines(x, dt(x, df = 9), lwd = 2)
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-167-1.pdf

2. The classical central limit theorem states that the standardized sum of i.i.d. random variables with finite variance converges in distribution to the standard normal distribution $N(0, 1)$. This exercise illustrates the central limit theorem.


- Write a simulation that performs the following steps:
- Generate a random sample X_1, \dots, X_5 of size $n = 5$ from the standard exponential distribution $Exp(1)$.
- Compute the sample sum.
- Repeat the steps $R = 10\,000$ times. For each replication, store the sum, e.g. into a vector Z .
- Plot the histogram of the sum and add the density function of $N(m, s^2)$ where m is the mean of Z and s is the standard deviation of Z .

```
clt_exp <- function(n) {
  R <- 10000
  Z <- rep(NA, R)
```

```

for (r in 1:R) {
  X <- rexp(n, rate = 1)
  Z[r] <- sum(X)
}
truehist(Z, col = "lightblue", main=paste("n =",n,sep=" "))
coord <- par("usr")
# par("usr") gives you a vector of the form c(x1, x2, y1, y2)
# giving the extremes of the coordinates of the plotting region
x <- seq(coord[1], coord[2], by = 0.1)
lines(x, dnorm(x, mean = mean(Z), sd = sd(Z)), lwd = 2)
}
clt_exp(5)


```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-168-1.pdf


- Increase the sample size n to $n = 50, 500, 5000$ and redo the exercise.

```
clt_exp(50)
```




IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-169-1.pdf

```
clt_exp(500)
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-169-2.pdf

```
clt_exp(5000)
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-169-3.pdf

- Redo the exercise with other distributions than the exponential. Use the uniform distribution, the t -distribution with 3 degrees of freedom, the Bernoulli distribution (i.e. binomial with parameter size=1), and the Poisson distribution.

```

clt <- function(n, distrib, df=3, lambda=5, prob=0.6) {
  R <- 10000
  Z <- rep(NA, R)
  for (r in 1:R) {

```

```


if (distrib == 1){
  X <- runif(n)
  strdist <- "Uniform"
}
if (distrib == 2){
  X <- rt(n, df = df)
  strdist <- "Student's t"
}
if (distrib == 3){
  X <- rbinom(n, size=1, prob=prob)
  strdist <- "Bernoulli"
}
if (distrib == 4){
  X <- rpois(n, lambda = lambda)
  strdist <- "Poisson"
}
Z[r] <- sum(X)
}
truehist(Z, col = "lightblue", xlab = strdist, main = paste("n =", n, sep = " "))
coord <- par("usr")
# par("usr") gives you a vector of the form c(x1, x2, y1, y2)
# giving the extremes of the coordinates of the plotting region
x <- seq(coord[1], coord[2], by = 0.1)
lines(x, dnorm(x, mean = mean(Z), sd = sd(Z)), lwd = 2)
}

```

```

par(mfrow = c(2,2))
for (n in c(5,50,500,5000)) {
  for (i in 1:4) {
    clt(n,i)
  }
}

```




- The central limit theorem breaks down if the variance of the summands is infinite. Redo the exercise using a t -distribution with only 1.5 degrees of freedom.

```

par(mfrow = c(1,1))
clt(500, 2, df = 1.5)

```



Linear regression

1. Load the Stata data set **wages.dta**. The variables are **earnings** (in Euro, 2009), **age**, **gender** (male=1, female=2), **education** (years of education), **hours** (hours worked during 2009), and **weight**.

```
library(foreign)
wages <- read.dta("data/wages.dta")

## Warning in read.dta("data/wages.dta"): value labels ('d1110211') for
## 'gender' are missing

head(wages)

##   gender age education hours earnings weight
## 1      2  51         18  1039      3900 737.73
## 2      2  59         18  2026     55550 459.86
## 3      1  22         13   312      2400 459.86
## 4      1  37         15  2338     54000 459.86
## 5      2  31         18  2858     28800 585.21
## 6      1  32         15  2078      9000 585.21

earnings <- wages$earnings
age <- wages$age
gender <- wages$gender
education <- wages$education
hours <- wages$hours
weight <- wages$weight
```

- Compute the (unweighted) wage equation

$$\ln \text{earnings}_i = \alpha + \beta_1 \text{age}_i + \beta_2 \text{age}_i^2 + \beta_3 \text{education}_i + \beta_4 \text{gender}_i + u_i,$$

print the summary of the `lm`-object, and interpret the output.

```
regr <- lm(log(earnings) ~ age + I(age^2) + education + gender)
summary(regr)

##
## Call:
## lm(formula = log(earnings) ~ age + I(age^2) + education + gender)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.4564 -0.3610  0.1617  0.5563  3.6001
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.903e+00  1.163e-01  50.76   <2e-16 ***
## age          1.663e-01  5.213e-03  31.91   <2e-16 ***
## I(age^2)     -1.726e-03  6.001e-05 -28.76   <2e-16 ***
## education    1.067e-01  3.043e-03  35.08   <2e-16 ***
## gender      -7.237e-01  1.660e-02 -43.60   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8947 on 11643 degrees of freedom
```

```
## Multiple R-squared:  0.2888, Adjusted R-squared:  0.2886
## F-statistic: 1182 on 4 and 11643 DF,  p-value: < 2.2e-16
```

- Add an interaction term for education and gender to the regression.

```
regr2 <- lm(log(earnings) ~ age + I(age^2) + education + gender + education:gender)
summary(regr2)
```

```
##
## Call:
## lm(formula = log(earnings) ~ age + I(age^2) + education + gender +
##     education:gender)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.4788 -0.3621  0.1572  0.5510  3.5676
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   6.610e+00  1.631e-01  40.527 < 2e-16 ***
## age           1.659e-01  5.205e-03  31.873 < 2e-16 ***
## I(age^2)      -1.717e-03  5.993e-05 -28.649 < 2e-16 ***
## education     5.139e-02  9.472e-03   5.426 5.88e-08 ***
## gender       -1.204e+00  7.954e-02 -15.134 < 2e-16 ***
## education:gender 3.763e-02  6.099e-03   6.170 7.05e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8933 on 11642 degrees of freedom
## Multiple R-squared:  0.2911, Adjusted R-squared:  0.2908
## F-statistic: 956.2 on 5 and 11642 DF,  p-value: < 2.2e-16
```

- Compute the weighted hourly wage equation

$$\ln \frac{\text{earnings}_i}{\text{hours}_i} = \alpha + \beta_1 \text{age}_i + \beta_2 \text{age}_i^2 + \beta_3 \text{education}_i + \beta_4 \text{gender}_i + u_i,$$


print the summary of the lm-object, and interpret the output.

```
regr3 <- lm(log(earnings/hours) ~ age + I(age^2) + education + gender, weights = weight)
summary(regr3)
```

```
##
## Call:
## lm(formula = log(earnings/hours) ~ age + I(age^2) + education +
##     gender, weights = weight)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -540.17   -9.18     0.00   14.33   367.44
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.929e-01  9.095e-02   5.419 6.11e-08 ***
## age           5.823e-02  4.010e-03  14.524 < 2e-16 ***
## I(age^2)      -5.494e-04  4.596e-05 -11.954 < 2e-16 ***
## education     8.069e-02  2.368e-03  34.072 < 2e-16 ***
```

```
## gender      -2.728e-01  1.264e-02 -21.581 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 36.63 on 10128 degrees of freedom
## Multiple R-squared:  0.1757, Adjusted R-squared:  0.1753
## F-statistic: 539.5 on 4 and 10128 DF,  p-value: < 2.2e-16
```

```
plot(regr3$residuals)
```



- Activate the packages `lmtest` and `sandwich`. Use the function `coeftest` to compute the heteroskedasticity robust standard errors (`vcov=vcovHC`) for the estimated coefficients.

```
library(lmtest)
```

```
## Loading required package: zoo
##
## Attaching package: 'zoo'
##
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
##
## Attaching package: 'lmtest'
##
## The following object is masked _by_ '.GlobalEnv':
##
##   wages
```

```
library(sandwich)
```

```
# Robust standard errors
```

```
coeftest(regr, vcov = vcovHC)
```

```
##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.9035e+00  1.4065e-01  41.974 < 2.2e-16 ***
## age          1.6634e-01  6.3384e-03  26.243 < 2.2e-16 ***
## I(age^2)     -1.7256e-03  7.2473e-05 -23.810 < 2.2e-16 ***
## education    1.0675e-01  2.9736e-03  35.898 < 2.2e-16 ***
## gender       -7.2369e-01  1.6586e-02 -43.633 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coeftest(regr2, vcov = vcovHC)
```

```
##
## t test of coefficients:
##
```

```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)   6.6102e+00 1.7589e-01 37.5822 < 2.2e-16 ***
## age          1.6591e-01 6.3388e-03 26.1741 < 2.2e-16 ***
## I(age^2)     -1.7168e-03 7.2509e-05 -23.6775 < 2.2e-16 ***
## education    5.1394e-02 8.8301e-03  5.8203 6.028e-09 ***
## gender       -1.2037e+00 7.9682e-02 -15.1059 < 2.2e-16 ***
## education:gender 3.7634e-02 6.0272e-03  6.2441 4.410e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coeftest(regr3, vcov = vcovHC)
```

```
## Warning in residuals^2/(1 - diag(hat))^2: Länge des längeren Objektes
##       ist kein Vielfaches der Länge des kürzeren Objektes
```

```
##
## t test of coefficients:
##
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.9288e-01 1.3111e-01  3.7592 0.0001714 ***
## age         5.8234e-02 6.1032e-03  9.5417 < 2.2e-16 ***
## I(age^2)    -5.4943e-04 7.1504e-05 -7.6838 1.687e-14 ***
## education   8.0690e-02 3.4333e-03 23.5025 < 2.2e-16 ***
## gender     -2.7276e-01 1.6960e-02 -16.0825 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Predict the hourly wage of a male person aged 60 years as a function of education (vary the years of education between 9 and 18). Set the option `se.fit=TRUE`. Inspect the object returned by the `predict` command. Plot the predicted values and add the ± 2 standard deviations confidence intervals.

```
forecast <- predict(regr3, newdata = data.frame(education = seq(9, 18, by = 0.5), age = 60, gender = 1)
names(forecast)
```

```
## [1] "fit"          "se.fit"       "df"           "residual.scale"
plot(seq(9, 18, by = 0.5), forecast$fit, type = "l", lwd = 2)
lines(seq(9, 18, by = 0.5), forecast$fit + 2 * forecast$se.fit, type = "l", col = "red", lwd = 1.5)
lines(seq(9, 18, by = 0.5), forecast$fit - 2 * forecast$se.fit, type = "l", col = "red", lwd = 1.5)
```

IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-181-1.pdf

2. Load the data set `bsp4.txt`.


```
bsp4 <- read.csv("data/bsp4.txt")
head(bsp4)
```

```
##      y      x
## 1 20.40 20.23
## 2 218.92 66.01
## 3 189.06 64.83
## 4 197.56 66.10
## 5 304.33 87.48
```

```
## 6 281.04 67.63
```

- Plot the scatter plot of y against x .


```
plot(bsp4$x, bsp4$y)
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-183-1.pdf

- Perform a simple linear regression of y on x and save the results as an `lm`-object `obj`. Add the regression line of y on x to the plot.

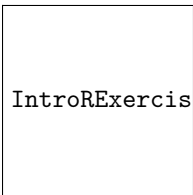
```
plot(bsp4$x, bsp4$y)
obj <- lm(bsp4$y ~ bsp4$x)
abline(obj)
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-184-1.pdf

- Extract the fitted values from `obj` and add them as red points to the plot (use the command `points`).

```
plot(bsp4$x, bsp4$y)
abline(obj)
points(bsp4$x, obj$fitted.values, col = "red")
```



IntroRExercisesWithSolutions_files/figure-latex/unnamed-chunk-185-1.pdf

- Extract the residuals of the regression and calculate the sum of the squared residuals, $SSR = \sum_{i=1}^{100} \hat{u}_i^2$

```
ssr <- sum((obj$residuals)^2)
print(ssr)
```

```
## [1] 223587.4
```

- Compute the total sum of squares and the explained sum of squares,

$$TSS = \sum_{i=1}^{100} (y_i - \bar{y})^2$$
$$ESS = \sum_{i=1}^{100} (\hat{y}_i - \bar{y})^2$$

and show that $ESS + SSR = TSS$.

```
tss <- sum((bsp4$y - mean(bsp4$y))^2)
ess <- sum((obj$fitted.values - mean(bsp4$y))^2)
ess + ssr - tss # this is numerically zero
```

```
## [1] 1.164153e-10
```

```
round(ess + ssr) == round(tss)
```

```
## [1] TRUE
```